



# KPLABS Course

HashiCorp Certified: Terraform Associate

Base Domain

**ISSUED BY**

Zeal

**REPRESENTATIVE**

[instructors@kplabs.in](mailto:instructors@kplabs.in)

# Base Domain - Getting Started & Setting Up Labs

## Module 1: Choosing a right Infrastructure as Code tool

### 1.1 Exploring Toolsets

There are various types of tools that can allow you to deploy infrastructure as code :

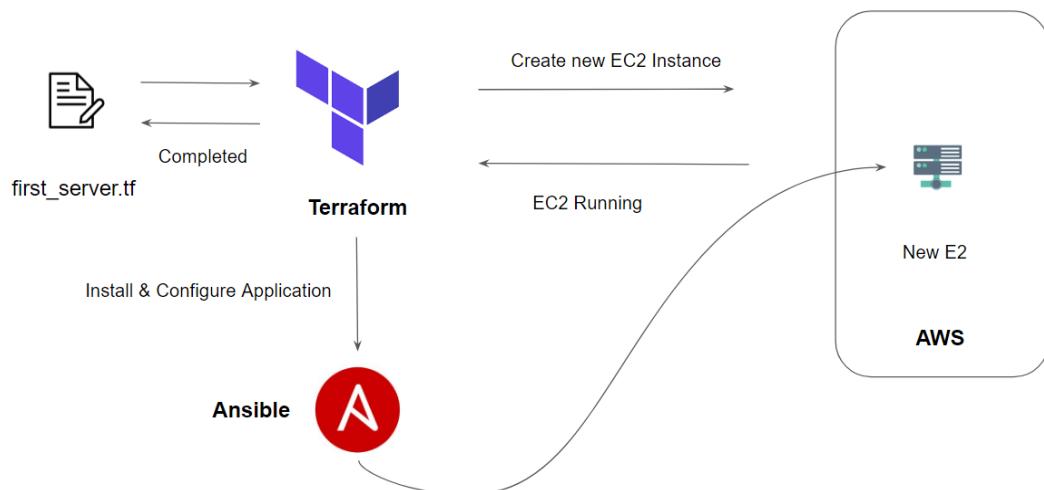
- Terraform
- CloudFormation
- Heat
- Ansible
- SaltStack
- Chef, Puppet and others

### 1.2 Configuration Management vs Infrastructure Orchestration

Ansible, Chef, Puppet are configuration management tools which means that they are primarily designed to install and manage software on existing servers.

Terraform, CloudFormation are the infrastructure orchestration tools which basically means they can provision the servers and infrastructure by themselves.

Configuration Management tools can do some degree of infrastructure provisioning, but the focus here is that some tools are going to be better fit for certain types of tasks.



### 1.3 Which tool to choose?

Question remains on how to choose right IAC tool for the organization

- i) Is your infrastructure going to be vendor specific in longer term ? Example AWS.
- ii) Are you planning to have multi-cloud / hybrid cloud based infrastructure ?
- iii) How well does it integrate with configuration management tools ?
- iv) Price and Support

### 1.4 Terraform

- i) Supports multiple platforms, has hundreds of providers.
- ii) Simple configuration language and faster learning curve.
- iii) Easy integration with configuration management tools like Ansible.
- iv) Easily extensible with the help of plugins.
- v) Free !!!

## **Module 2: Installation Process of Terraform**

Terraform installation is very simple.

You have a single binary file, download and use it.



Terraform works on multiple platforms, these includes:

- Windows
- macOS
- Linux
- FreeBSD
- OpenBSD
- Solaris

You can download Terraform binary through the following link:

<https://www.terraform.io/downloads>

## Module 3: Setting up the Lab

i) Create a new AWS Account.

ii) Begin the course



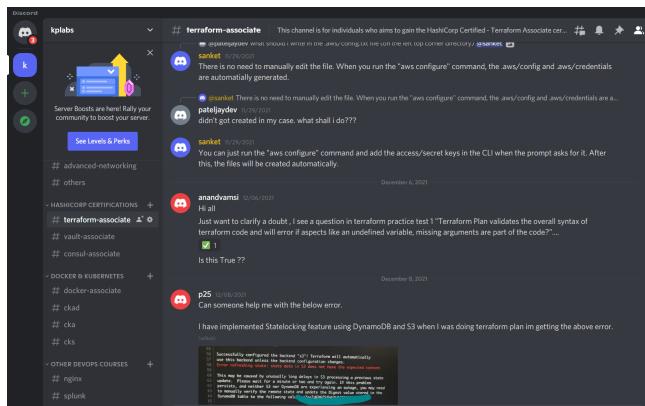
The screenshot shows the AWS Free Tier landing page. At the top, there is a navigation bar with links for 'Menu', 'aws' logo, 'Contact Sales', 'Products', 'Solutions', 'Pricing', 'More', 'English', 'My Account', and a yellow 'Create an AWS Account' button. The main header is 'AWS Free Tier'. Below it, a sub-header reads: 'The AWS Free Tier enables you to gain free, hands-on experience with the AWS platform, products, and services.' A large yellow 'Create a Free Account' button is centered. At the bottom of the main section, there are three links: 'Free Tier Details', 'Get Started', and 'Free Tier Software'. The 'Free Tier Details' section is expanded, showing a sub-header 'AWS Free Tier Details' and several filter options: '★ FEATURED', '12 MONTHS FREE', 'ALWAYS FREE', 'TRIALS', 'PRODUCT CATEGORIES', and 'ALL'.

# Join Our Discord Community

We invite you to join our Discord community, where you can interact with our support team for any course-based technical queries and connect with other students who are doing the same course.

Joining URL:

<http://kplabs.in/chat>





# KPLABS Course

HashiCorp Certified: Terraform Associate

## Domain 1

**ISSUED BY**

Zeal

**REPRESENTATIVE**

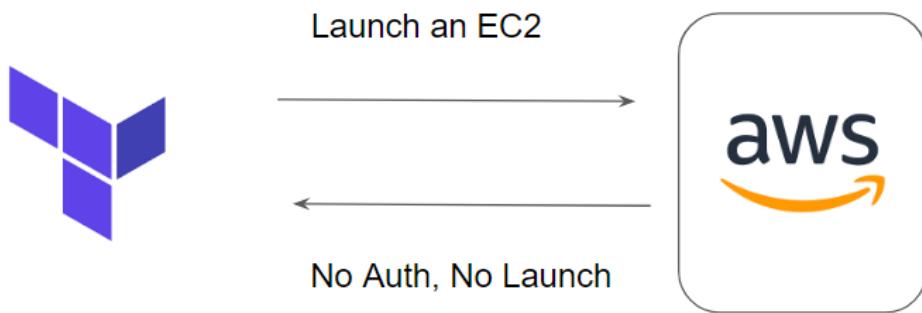
[instructors@kplabs.in](mailto:instructors@kplabs.in)

# Domain 1 - Deploying Infrastructure with Terraform

## Module 1: Creating first EC2 instance with Terraform

When you want to launch a resource in AWS, there are three important considerations.

- How will you authenticate to AWS?
- Which region the resource needs to be launched.
- Which resource you want to launch.



Each Provider has resources divided into multiple categories.

Resources require a certain set of arguments based on the use-case.

## Module 2: Provider and Resources

### 2.1 Overview of Providers:

Terraform supports multiple providers.

Depending on what type of infrastructure we want to launch, we have to use appropriate providers accordingly.



### 2.2 Initialization Phase

Upon adding a provider, it is important to run `terraform init` which in-turn will download plugins associated with the provider.

```
C:\Users\Zeal Vora\Desktop\kplabs-terraform>terraform init
Initializing the backend...
Initializing provider plugins...
- Reusing previous version of hashicorp/aws from the dependency lock file
- Finding latest version of hashicorp/azurerm...
- Installing hashicorp/aws v3.26.0...
- Installed hashicorp/aws v3.26.0 (signed by HashiCorp)
- Installing hashicorp/azurerm v2.45.1...
- Installed hashicorp/azurerm v2.45.1 (signed by HashiCorp)

Terraform has made some changes to the provider dependency selections recorded
in the .terraform.lock.hcl file. Review those changes and commit them to your
version control system if they represent changes you intended to make.

Terraform has been successfully initialized!
```

## 2.3 Resources

Resources are the reference to the individual services which the provider has to offer

Example:

- resource aws\_instance
- resource aws\_alb
- resource iam\_user
- resource digitalocean\_droplet

```
resource "aws_instance" "myec2" {  
    ami = "ami-082b5a644766e0e6f"  
    instance_type = "t2.micro"  
}
```

## 2.4 Important Update - Newer Version

From 0.13 onwards, Terraform requires explicit source information for any providers that are not HashiCorp-maintained, using a new syntax in the required\_providers nested block inside the Terraform configuration block

```
provider "aws" {
  region      = "us-west-2"
  access_key  = "PUT-YOUR-ACCESS-KEY-HERE"
  secret_key  = "PUT-YOUR-SECRET-KEY-HERE"
}
```

HashiCorp Maintained

```
terraform {
  required_providers {
    digitalocean = {
      source = "digitalocean/digitalocean"
    }
  }
}

provider "digitalocean" {
  token = "PUT-YOUR-TOKEN-HERE"
}
```

Non-HashiCorp Maintained



```
provider "aws" {
  region      = "us-west-2"
  access_key  = "PUT-YOUR-ACCESS-KEY-HERE"
  secret_key  = "PUT-YOUR-SECRET-KEY-HERE"
}
```



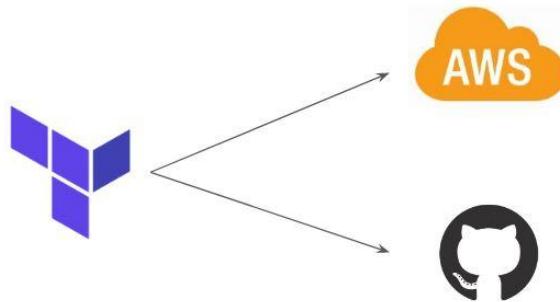
```
terraform {
  required_providers {
    digitalocean = {
      source = "digitalocean/digitalocean"
    }
  }
}

provider "digitalocean" {
  token = "PUT-YOUR-TOKEN-HERE"
}
```

Terraform 0.13 onwards

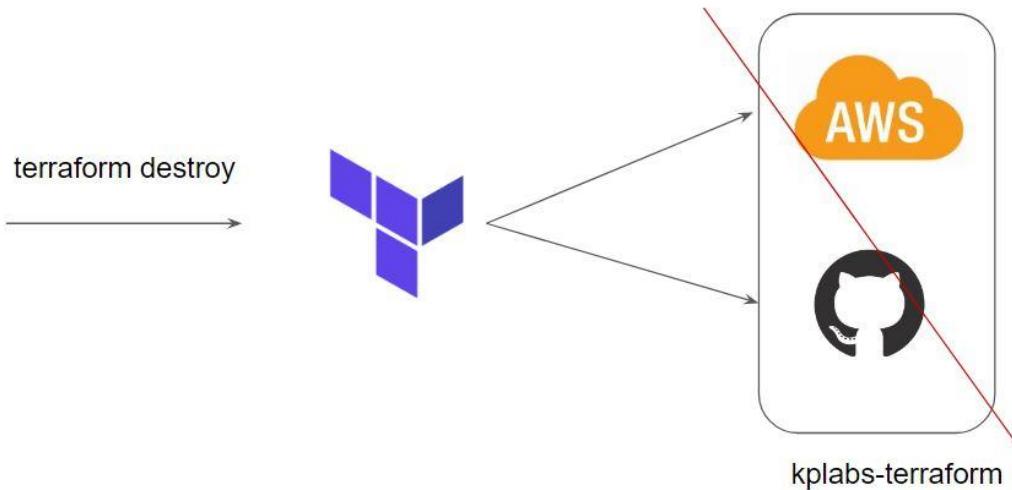
## Module 3: Destroying Infrastructure with Terraform (NEW)

If you keep the infrastructure running, you will get charged for it. Hence it is important for us to also know how we can delete the infrastructure resources created via terraform.



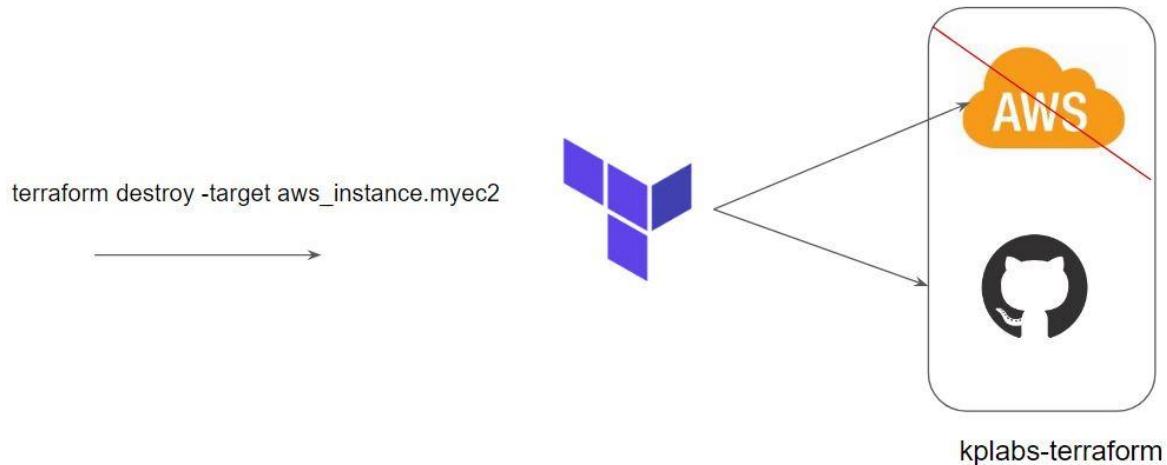
### 3.1 Approach 1

terraform destroy allows us to destroy all the resources that are created within the folder.



### 3.2 Approach 2

terraform destroy with -target flag allows us to destroy the specific resource.



### 3.3 Terraform Destroy with Target

The -target option can be used to focus Terraform's attention on only a subset of resources.

Combination of: Resource Type + Local Resource Name

Resource Type	Local Resource Name
aws_instance	myec2
github_repository	example

```
resource "aws_instance" "myec2" {
    ami = "ami-082b5a644766e0e6f"
    instance_type = "t2.micro"
}
```

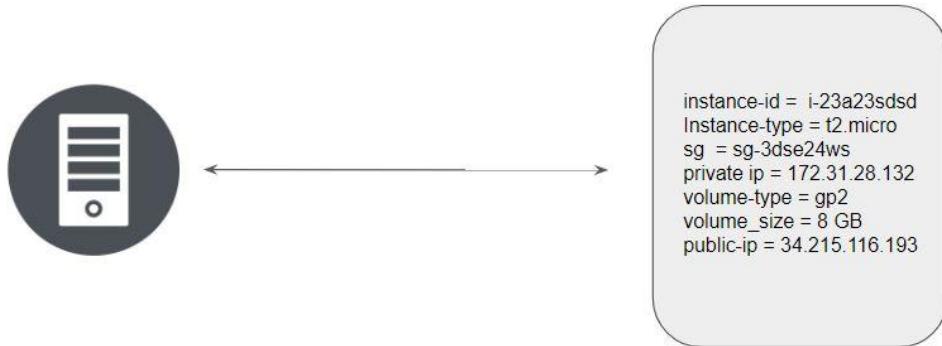
```
resource "github_repository" "example" {
    name      = "terraform-repo"
    visibility = "private"
}
```

## Module 4: Terraform State File

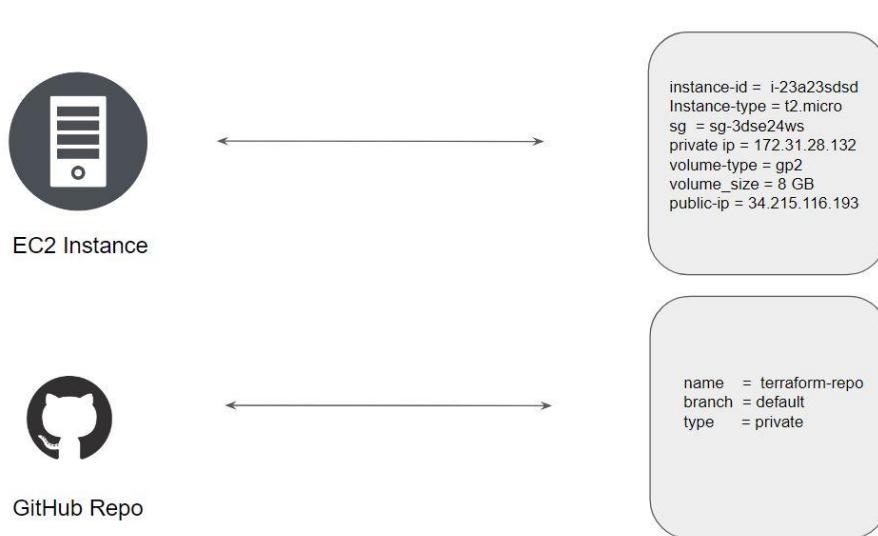
### 4.1 Overview of State Files:

Terraform stores the state of the infrastructure that is being created from the TF files.

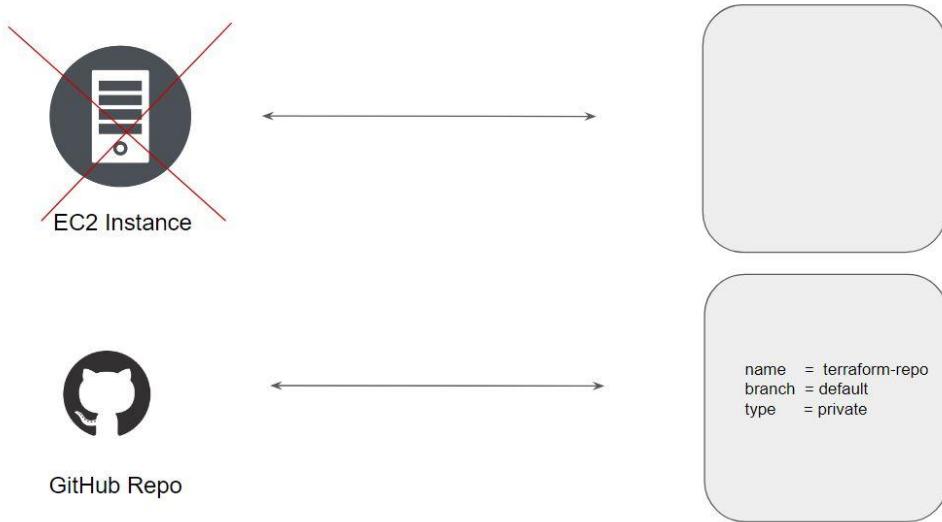
This state allows terraform to map real-world resources to your existing configuration.



Multiple resources in Terraform will have a separate block with the state file.



Whenever a resource is removed, its corresponding entry under the state file is also removed.



## Module 5: Desired & Current State

### 5.1 Desired State

Terraform's primary function is to create, modify, and destroy infrastructure resources to match the desired state described in a Terraform configuration

```
resource "aws_instance" "myec2" {
    ami = "ami-082b5a644766e0e6f"
    instance_type = "t2.micro"
}
```

### 5.2 Current State

The current state is the actual state of a resource that is currently deployed.

```
resource "aws_instance" "myec2" {
    ami = "ami-082b5a644766e0e6f"
    instance_type = "t2.micro"
}
```



t2.medium

### 5.3 Important Pointer

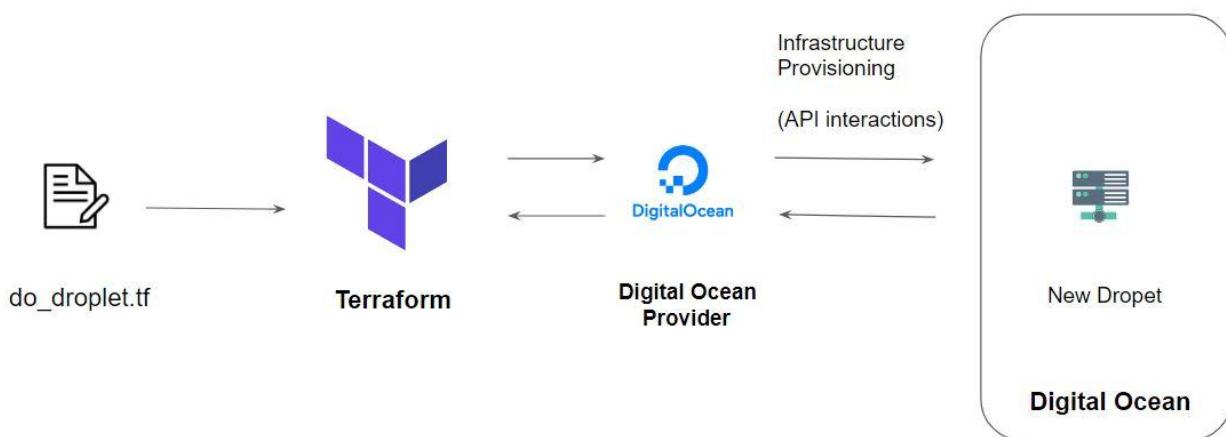
Terraform tries to ensure that the deployed infrastructure is based on the desired state.

If there is a difference between the two, terraform plan presents a description of the changes necessary to achieve the desired state.



## Module 6: Provider Versioning

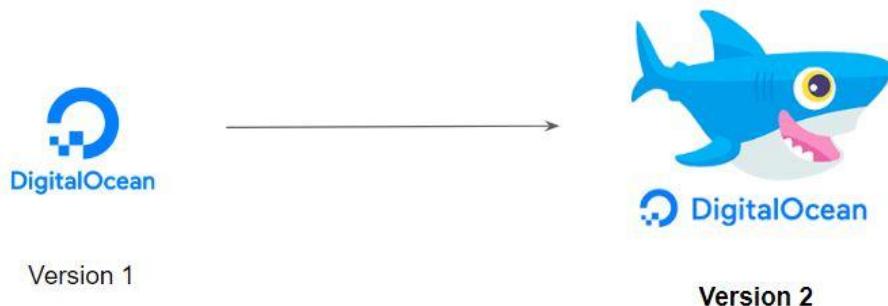
### 6.1 Overview of Provider Architecture



### 6.2 Provider Versioning

Provider plugins are released separately from Terraform itself.

They have a different set of version numbers.



### 6.3 Explicitly Setting Provider Version

During terraform init, if the version argument is not specified, the most recent provider will be downloaded during initialization.

For production use, you should constrain the acceptable provider versions via configuration, to ensure that new versions with breaking changes will not be automatically installed

```
provider "aws" {
    region      = "us-west-2"
    version     = "2.7"
}
```

### 6.4 Arguments for Specifying the provider

There are multiple ways of specifying the version of a provider.

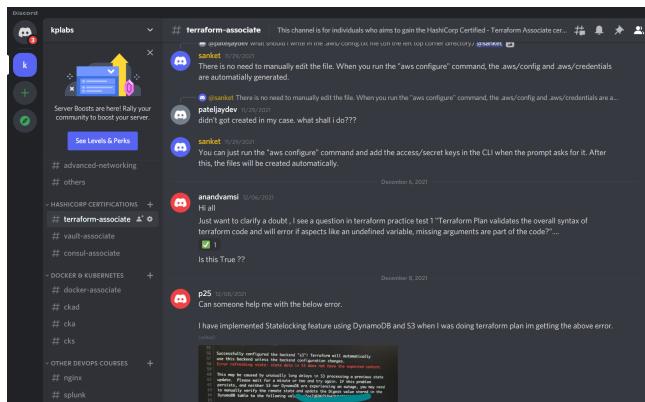
Version Number Arguments	Description
$\geq 1.0$	Greater than equal to the version
$\leq 1.0$	Less than equal to the version
$\sim >2.0$	Any version in the 2.X range.
$\geq 2.10, \leq 2.30$	Any version between 2.10 and 2.30

## Join Our Discord Community

We invite you to join our Discord community, where you can interact with our support team for any course-based technical queries and connect with other students who are doing the same course.

Joining URL:

<http://kplabs.in/chat>







# KPLABS Course

HashiCorp Certified: Terraform Associate

## Domain 2

**ISSUED BY**

Zeal

**REPRESENTATIVE**

[instructors@kplabs.in](mailto:instructors@kplabs.in)

## Domain 2 - Read, Generate, Modify Configurations

### Module 1: Attributes and Output Values

Terraform has the capability to output the attribute of a resource with the output values.

Example:

- ec2\_public\_ip = 35.161.21.197
- bucket\_identifier = terraform-test-kplabs.s3.amazonaws.com

An outputted attributes can not only be used for the user reference but it can also act as an input to other resources being created via terraform

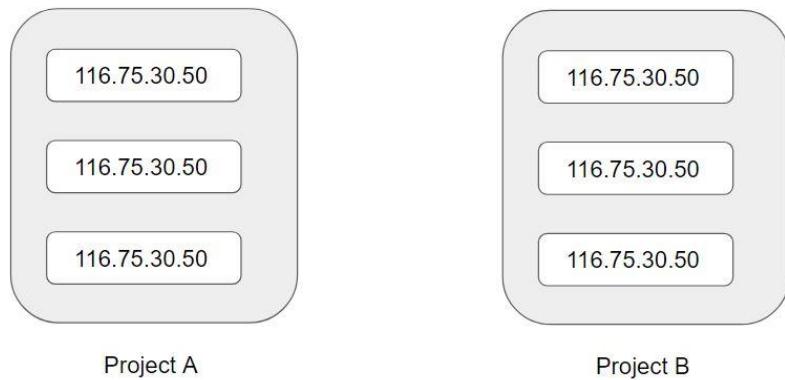
Let's understand this with an example:

After EIP gets created, its IP address should automatically get whitelisted in the security group.

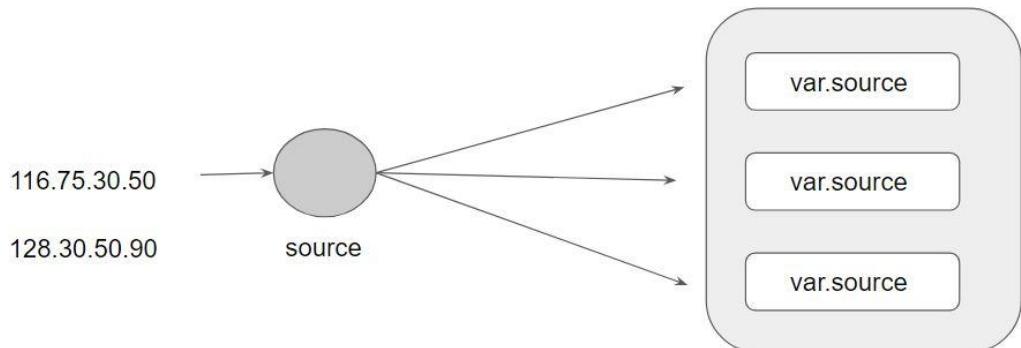


## Module 2: Terraform variables

Repeated static values can create more work in the future.



Terraform Variables allows us to centrally define the values that can be used in multiple terraform configuration blocks.



## Module 3: Approach for Variable Assignment

Variables in Terraform can be assigned values in multiple ways.

Some of these include:

- Environment variables
- Command Line Flags
- From a File
- Variable Defaults

### Sample Commands for the following:

i) Environment Variables:

```
export TF_VAR_instancetype="t2.nano"
echo $TF_VAR
```

ii) Command Line Flags:

```
terraform plan -var="instancetype=t2.small"
terraform plan -var-file="custom.tfvars"
```

iii) From a File (terraform.tfvars):

```
instancetype="t2.large"
```

iv) Variable Defaults:

```
variable "instancetype" {
  default = "t2.micro"
}
```

## Module 4: Data Types for Variables

### 4.1 Overview of Type Constraints

The type argument in a variable block allows you to restrict the type of value that will be accepted as the value for a variable

```
variable "image_id" {  
    type = string  
}
```

If no type constraint is set then a value of any type is accepted.

### 4.2 Example Use-Case 1

Every employee in Medium Corp is assigned an Identification Number.

Any resource that an employee creates should be created with the name of the identification number only.

variables.tf	terraform.tfvars
variable "instance_name" {}	instance_name="john-123"

### 4.3 Example Use-Case 2

Every employee in Medium Corp is assigned an Identification Number.

Any EC2 instance that employee creates should be created using the identification number only.

variables.tf	terraform.tfvars
variable "instance_name" { type=number }	instance_name="john-123"

## 4.4 Overview of Data Types

Type Keywords	Description
string	Sequence of Unicode characters representing some text, like "hello".
list	Sequential list of values identified by their position. Starts with 0 ["mumbai", "singapore", "usa"]
map	a group of values identified by named labels, like {name = "Mabel", age = 52}.
number	Example: 200

## Module 5: Count Parameter

### 5.1 Overview of Count:

The count parameter on resources can simplify configurations and let you scale resources by simply incrementing a number.

Let's assume, you need to create two EC2 instances. One of the common approaches is to define two separate resource blocks for aws\_instance.

```
resource "aws_instance" "instance-1" {  
    ami = "ami-082b5a644766e0e6f"  
    instance_type = "t2.micro"  
}
```



```
resource "aws_instance" "instance-2" {  
    ami = "ami-082b5a644766e0e6f"  
    instance_type = "t2.micro"  
}
```

With the count parameter, we can simply specify the count value and the resource can be scaled accordingly.

```
resource "aws_instance" "instance-1" {
    ami = "ami-082b5a644766e0e6f"
    instance_type = "t2.micro"
    count = 5
}
```

## 5.2 Count Index

In resource blocks where the count is set, an additional count object is available in expressions, so you can modify the configuration of each instance.

This object has one attribute:

count.index — The distinct index number (starting with 0) corresponding to this instance.

## 5.3 Challenges with Count Parameter

With the below code, terraform will create 5 IAM users. But the problem is that all will have the same name.

```
resource "aws_iam_user" "lb" {
    name = "loadbalancer"
    count = 5
    path = "/system/"
}
```

count.index allows us to fetch the index of each iteration in the loop.

```
resource "aws_iam_user" "lb" {
    name = "loadbalancer.${count.index}"
    count = 5
    path = "/system/"
}
```

## Understanding Challenge with Default Count Index

Having a username like loadbalancer0, loadbalancer1 might not always be suitable.

Better names like dev-loadbalancer, stage-loadbalancer, prod-loadbalancer is better.

count.index can help in such a scenario as well.

```
variable "elb_names" {
  type      = list
  default   = ["dev-loadbalancer", "stage-loadbalancer", "prod-loadbalancer"]
}
```

## Module 6: Conditional Expression

A conditional expression uses the value of a bool expression to select one of two values.

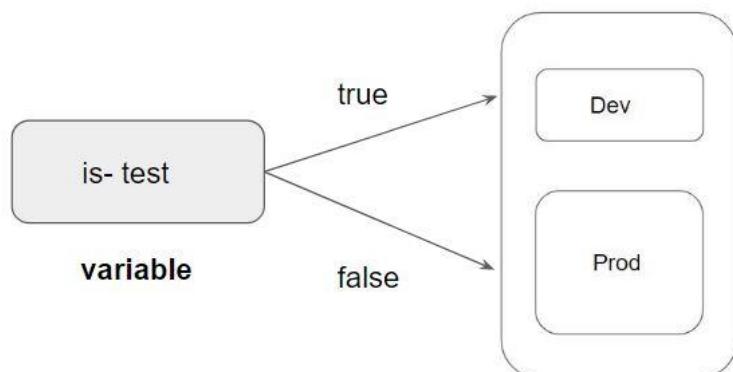
Syntax of Conditional expression:

```
condition ? true_val : false_val
```

If the condition is true then the result is true\_val. If the condition is false then the result is false\_val.

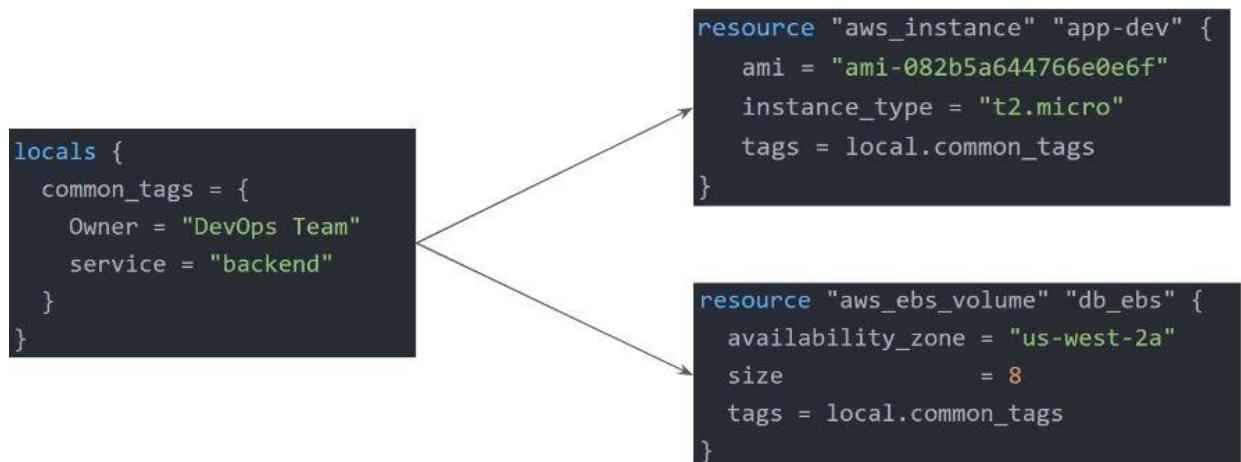
Let's assume that there are two resource blocks as part of terraform configuration.

Depending on the variable value, one of the resource blocks will run.



## Module 7: Local Values

A local value assigns a name to an expression, allowing it to be used multiple times within a module without repeating it.



### Local Values Support for Expression

Local Values can be used for multiple different use-cases like having a conditional expression.

```
locals {
    name_prefix = "${var.name != "" ? var.name : var.default}"
}
```

### Important Pointers for Local Values:

Local values can be helpful to avoid repeating the same values or expressions multiple times in a configuration.

If overused they can also make a configuration hard to read by future maintainers by hiding the actual values used

Use local values only in moderation, in situations where a single value or result is used in many places and that value is likely to be changed in the future.

## Module 8: Terraform Functions

The Terraform language includes a number of built-in functions that you can use to transform and combine values.

The general syntax for function calls is a function name followed by comma-separated arguments in parentheses:

```
function (argument1, argument2)
```

Example:

```
> max(5, 12, 9)
```

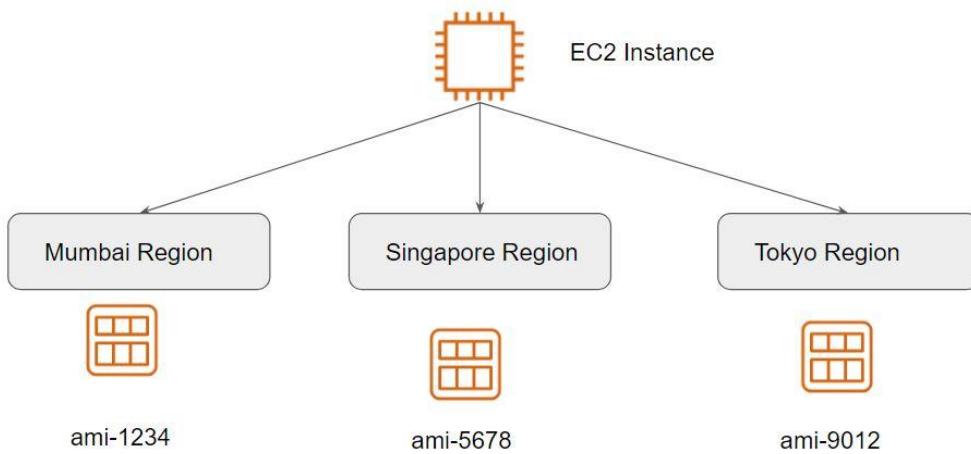
```
12
```

The Terraform language does not support user-defined functions, and so only the functions built into the language are available for use

- Numeric
- String
- Collection
- Encoding
- Filesystem
- Date and Time
- Hash and Crypto
- IP Network
- Type Conversion

## Module 9: Data Sources

Data sources allow data to be fetched or computed for use elsewhere in Terraform configuration.



A data source is defined under the data block.

It reads from a specific data source (aws\_ami) and exports results under “app\_ami”

```

data "aws_ami" "app_ami" {
  most_recent = true
  owners      = ["amazon"]

  filter {
    name   = "name"
    values = ["amzn2-ami-hvm*"]
  }
}

resource "aws_instance" "instance-1" {
  ami     = data.aws_ami.app_ami.id
  instance_type = "t2.micro"
}

```

## Module 10: Debugging in Terraform

Terraform has detailed logs that can be enabled by setting the TF\_LOG environment variable to any value.

You can set TF\_LOG to one of the log levels TRACE, DEBUG, INFO, WARN or ERROR to change the verbosity of the logs

```
bash-4.2# terraform plan
2020/04/22 13:45:31 [INFO] Terraform version: 0.12.24
2020/04/22 13:45:31 [INFO] Go runtime version: go1.12.13
2020/04/22 13:45:31 [INFO] CLI args: []string{"/usr/bin/terraform", "plan"}
2020/04/22 13:45:31 [DEBUG] Attempting to open CLI config file: /root/.terraformrc
2020/04/22 13:45:31 [DEBUG] File doesn't exist, but doesn't need to. Ignoring.
2020/04/22 13:45:31 [DEBUG] checking for credentials in "/root/.terraform.d/plugins"
2020/04/22 13:45:31 [INFO] CLI command args: []string{"plan"}
2020/04/22 13:45:31 [TRACE] Meta.Backend: built configuration for "s3" backend with hash value 789489680
2020/04/22 13:45:31 [TRACE] Meta.Backend: backend has not previously been initialized in this working directory
2020/04/22 13:45:31 [DEBUG] New state was assigned lineage "a10f92bf-686d-e6cf-3e9d-755be5c8a6a3"
2020/04/22 13:45:31 [TRACE] Meta.Backend: moving from default local state only to "s3" backend
```

### Important Pointers for Debugging:

TRACE is the most verbose and it is the default if TF\_LOG is set to something other than a log level name.

To persist logged output you can set TF\_LOG\_PATH in order to force the log to always be appended to a specific file when logging is enabled.

## Module 11: Terraform Format

Anyone who is into programming knows the importance of formatting the code for readability.

```
provider "aws" {
    region      = "us-west-2"
    access_key  = "AKIAQIW66DN2W7WOYRGY"
    secret_key  = "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"
    version     = ">=2.10,<=2.30"
}
```

The terraform fmt command is used to rewrite Terraform configuration files to take care of the overall formatting

```

provider "aws" {
    region      = "us-west-2"
    access_key  = "AKIAQIW66DN2W7WOYRGY"
    secret_key  = "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"
    version     = ">=2.10,<=2.30"
}

↓

provider "aws" {
    region      = "us-west-2"
    access_key  = "AKIAQIW66DN2W7WOYRGY"
    secret_key  = "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"
    version     = ">=2.10,<=2.30"
}

```

**Before fmt**

**After fmt**

## Module 12: Terraform Validate

Terraform Validate primarily checks whether a configuration is syntactically valid.

It can check various aspects including unsupported arguments, undeclared variables, and others.

```

resource "aws_instance" "myec2" {
    ami          = "ami-082b5a644766e0e6f"
    instance_type = "t2.micro"
    sky          = "blue"
}

→ bash-4.2# terraform validate
      Error: Unsupported argument
      on validate.tf line 10, in resource "aws_instance" "myec2":
10:   sky = "blue"

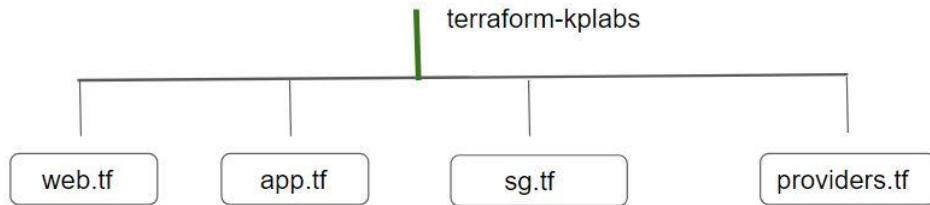
      An argument named "sky" is not expected here.

```

## Module 13: Load Order & Semantics

Terraform generally loads all the configuration files within the directory specified in alphabetical order.

The files loaded must end in either .tf or .tf.json to specify the format that is in use.



## Module 14: Dynamic Blocks

### 14.1 Understanding the Challenge:

In many of the use-cases, there are repeatable nested blocks that need to be defined.

This can lead to a long code and it can be difficult to manage in a long time.

```
ingress {  
    from_port  = 9200  
    to_port    = 9200  
    protocol   = "tcp"  
    cidr_blocks = ["0.0.0.0/0"]  
}
```

```
ingress {  
    from_port  = 8300  
    to_port    = 8300  
    protocol   = "tcp"  
    cidr_blocks = ["0.0.0.0/0"]  
}
```

### 14.2 Overview of Dynamic Blocks

Dynamic Block allows us to dynamically construct repeatable nested blocks which is supported inside resource, data, provider, and provisioner blocks:

```
dynamic "ingress" {  
    for_each = var.ingress_ports  
    content {  
        from_port  = ingress.value  
        to_port    = ingress.value  
        protocol   = "tcp"  
        cidr_blocks = ["0.0.0.0/0"]  
    }  
}
```

## 14.3 Overview of Iterators

The iterator argument (optional) sets the name of a temporary variable that represents the current element of the complex value

If omitted, the name of the variable defaults to the label of the dynamic block ("ingress" in the example above).

```
dynamic "ingress" {
  for_each = var.ingress_ports
  content {
    from_port    = ingress.value
    to_port      = ingress.value
    protocol     = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```



```
dynamic "ingress" {
  for_each = var.ingress_ports
  iterator = port
  content {
    from_port    = port.value
    to_port      = port.value
    protocol     = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }
}
```

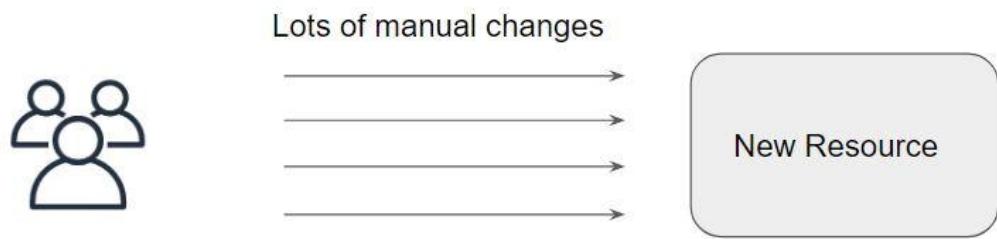
## **Module 15: Terraform Taint**

### 15.1 Understanding the Challenge:

You have created a new resource via Terraform.

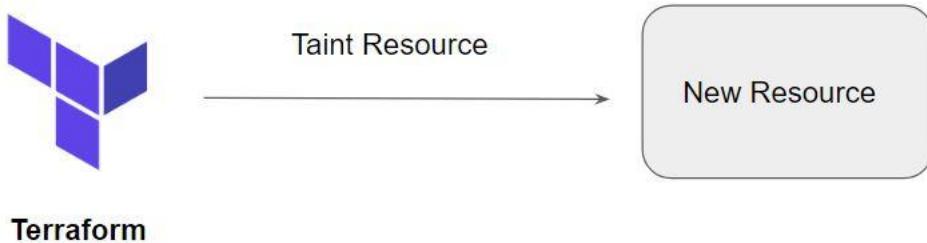
Users have made a lot of manual changes (both infrastructure and inside the server)

Two ways to deal with this: Import The Changes to Terraform / Delete & Recreate the resource



## 15.2 Overview of Terraform Taint

The `terraform taint` command manually marks a Terraform-managed resource as tainted, forcing it to be destroyed and recreated on the next apply.



## 15.3 Important Pointers for Terraform Taint

This command will not modify infrastructure but does modify the state file in order to mark a resource as tainted.

Once a resource is marked as tainted, the next plan will show that the resource will be destroyed and recreated and the next apply will implement this change.

Note that tainting a resource for recreation may affect resources that depend on the newly tainted resource.

## Module 16: Splat Expression

Splat Expression allows us to get a list of all the attributes.

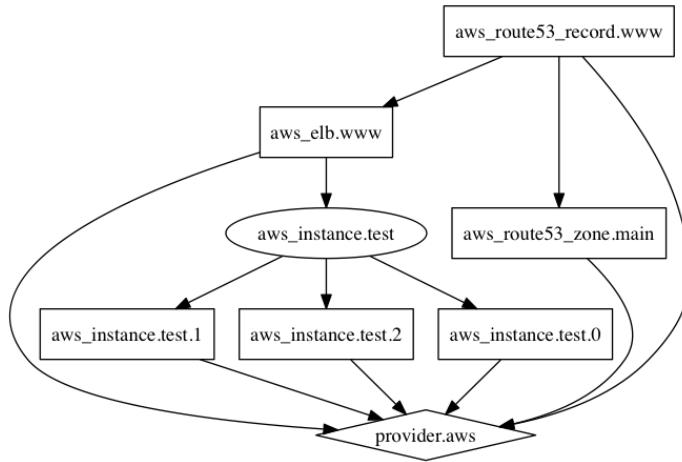
```
resource "aws_iam_user" "lb" {
    name = "iamuser.${count.index}"
    count = 3
    path = "/system/"
}

output "arns" {
    value = aws_iam_user.lb[*].arn
}
```

## Module 17: Terraform Graph

The terraform graph command is used to generate a visual representation of either a configuration or execution plan

The output of [terraform graph](#) is in the DOT format, which can easily be converted to an image.



## Module 18: Saving Terraform Plan to a File

The generated terraform plan can be saved to a specific path.

This plan can then be used with `terraform apply` to be certain that only the changes shown in this plan are applied.

Example:

```
terraform plan -out=path
```

## Module 19: Terraform Output

The `terraform output` command is used to extract the value of an output variable from the state file.

```
C:\Users\Zeal Vora\Desktop\terraform>terraform output iam_names
[
  "iamuser.0",
  "iamuser.1",
  "iamuser.2",
]
```

## Module 20: Terraform Settings

The special terraform configuration block type is used to configure some behaviors of Terraform itself, such as requiring a minimum Terraform version to apply your configuration.

Terraform settings are gathered together into terraform blocks:

```
terraform {
  # ...
}
```

### 20.1 Setting 1 - Terraform Version

The [required\\_version](#) setting accepts a version constraint string, which specifies which versions of Terraform can be used with your configuration.

If the running version of Terraform doesn't match the constraints specified, Terraform will produce an error and exit without taking any further actions.

```
terraform {
  required_version = "> 0.12.0"
}
```

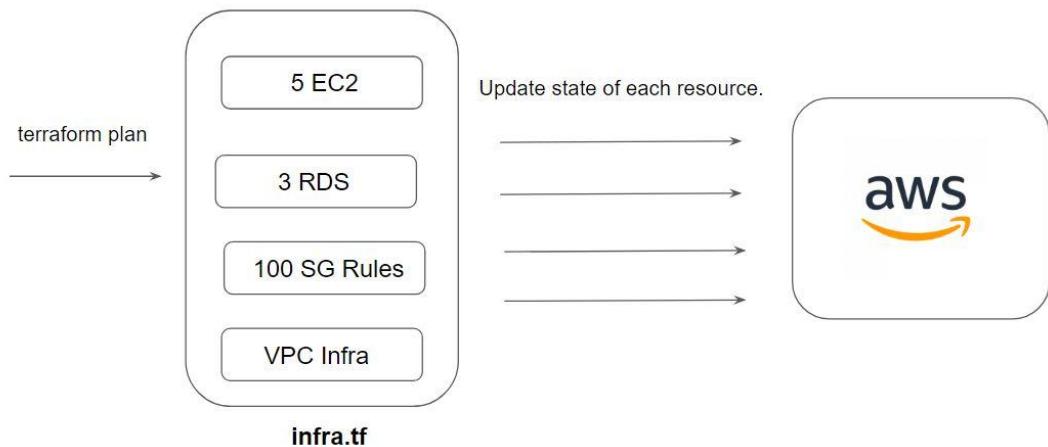
## 20.1 Setting 2 - Provider Version

The `required_providers` block specifies all of the providers required by the current module, mapping each local provider name to a source address and a version constraint.

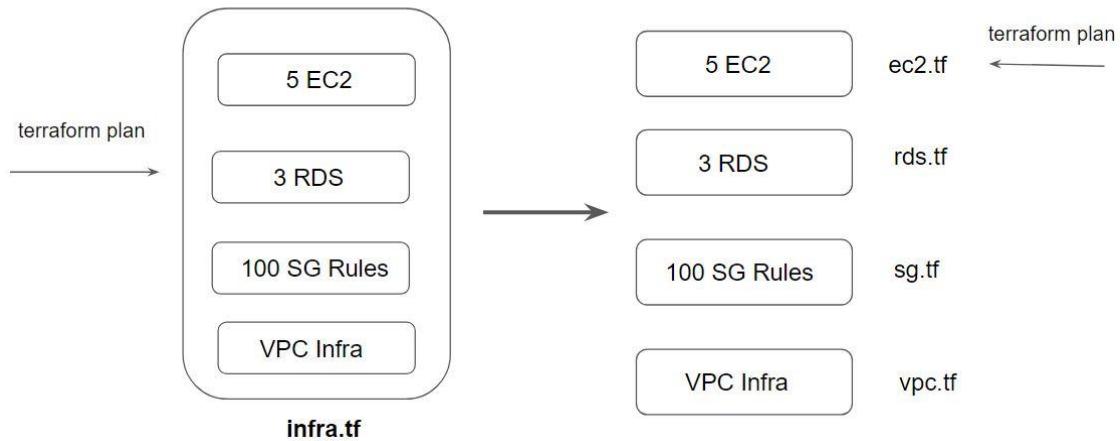
```
terraform {
  required_providers {
    mycloud = {
      source  = "mycorp/mycloud"
      version = "~> 1.0"
    }
  }
}
```

## Module 21: Dealing with Large Infrastructure

When you have a larger infrastructure, you will face issues related to API limits for a provider.



It is important to switch to a smaller configurations were each can be applied independently.



## 21.1 Setting Refresh to False

We can prevent terraform from querying the current state during operations like terraform plan.

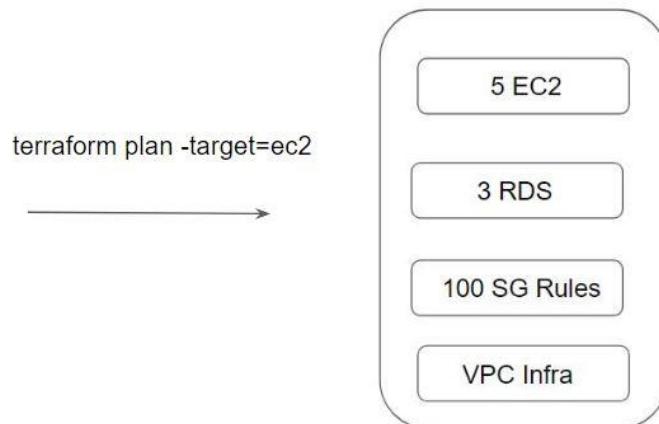
This can be achieved with the [-refresh=false](#) flag



## 21.2 Specify the Target

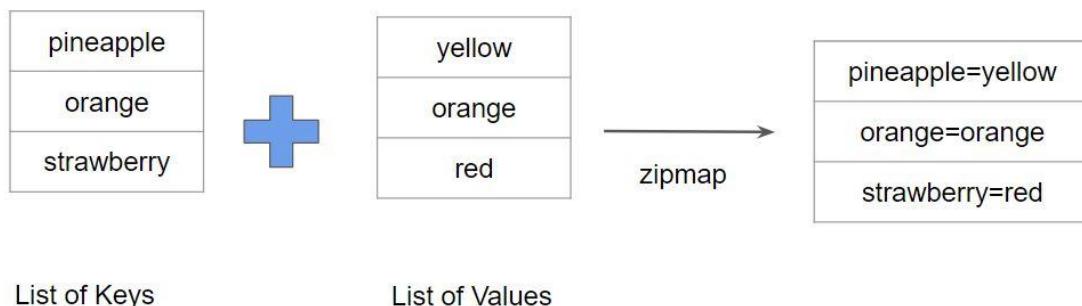
The `-target=resource` flag can be used to target a specific resource.

Generally used as a means to operate on isolated portions of very large configurations



## Module 22: Zipmap Function

The zipmap function constructs a map from a list of keys and a corresponding list of values.



Following screenshot shows a sample output of Zipmap

```
←[J> zipmap(["pineapple", "oranges", "strawberry"], ["yellow", "orange", "red"])
{
    "oranges" = "orange"
    "pineapple" = "yellow"
    "strawberry" = "red"
}
```

### Let us understand Zipmap with a sample use-case

You are creating multiple IAM users.

You need output which contains direct mapping of IAM names and ARNs

```
zipmap = {
    "demo-user.0" = "arn:aws:iam::018721151861:user/system/demo-user.0"
    "demo-user.1" = "arn:aws:iam::018721151861:user/system/demo-user.1"
    "demo-user.2" = "arn:aws:iam::018721151861:user/system/demo-user.2"
}
```

## Module 23: Comments in Terraform Code

A comment is a text note added to source code to provide explanatory information, usually about the function of the code

```
'''In this program, we check if the number is positive or
negative or zero and
display an appropriate message'''

num = 3.4

# Try these two variations as well:
# num = 0
# num = -4.5

if num > 0:
    print("Positive number")
elif num == 0:
    print("Zero")
else:
    print("Negative number")
```

The Terraform language supports three different syntaxes for comments:

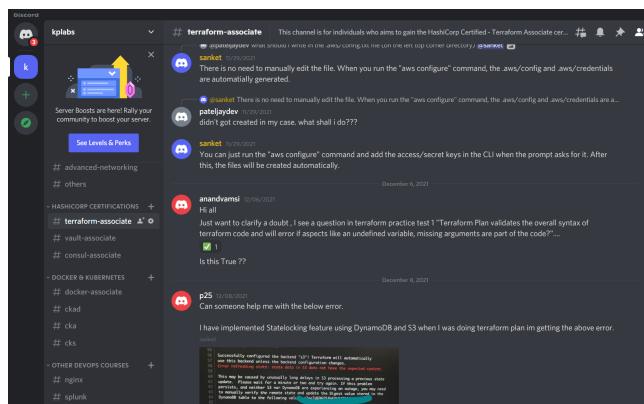
Type	Description
#	begins a single-line comment, ending at the end of the line.
//	also begins a single-line comment, as an alternative to #.
/* and */	are start and end delimiters for a comment that might span over multiple lines.

## Join Our Discord Community

We invite you to join our Discord community, where you can interact with our support team for any course-based technical queries and connect with other students who are doing the same course.

Joining URL:

<http://kplabs.in/chat>





# KPLABS Course

HashiCorp Certified: Terraform Associate

## Domain 3 - Terraform Provisioners

**ISSUED BY**

Zeal

**REPRESENTATIVE**

[instructors@kplabs.in](mailto:instructors@kplabs.in)

## Module 1: Understanding Provisioners in Terraform

### 1.1 Understanding the Challenge

Till now we have been working only on the creation and destruction of infrastructure scenarios.

Let's take an example:

We created a web-server EC2 instance with Terraform.

Problem: It is only an EC2 instance, it does not have any software installed.

What if we want a complete end to end solution?

### 1.2 Introducing Terraform Provisioners

Provisioners are used to execute scripts on a local or remote machine as part of resource creation or destruction.

Let's take an example:

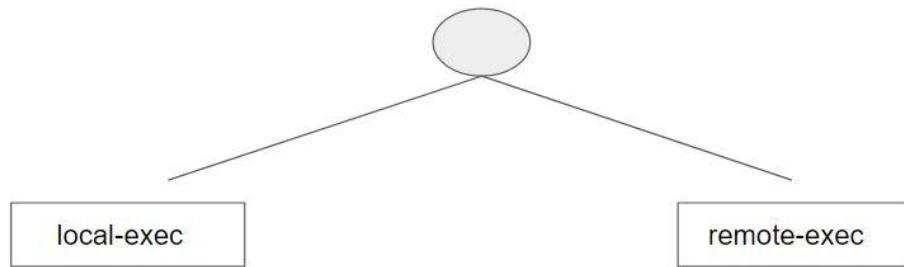
On creation of Web-Server, execute a script which installs Nginx web-server.



## Module 2: Types of Provisioners

Terraform has the capability to turn provisioners both at the time of resource creation as well as destruction.

There are two main types of provisioners:



### 2.1 Local Exec Provisioners

local-exec provisioners allow us to invoke a local executable after the resource is created.

One of the most used approaches of local-exec is to run ansible-playbooks on the created server after the resource is created.

Let's take an example:

```
provisioner "local-exec" {
  command = "echo ${aws_instance.web.private_ip} >> private_ips.txt"
}
```

## 2.2 Remote Exec Provisioners

Remote-exec provisioners allow invoking scripts directly on the remote server.

Let's take an example:

```
resource "aws_instance" "web" {  
  # ...  
  
  provisioner "remote-exec" {  
    .....  
  }  
}
```

## **Module 3: Provisioner Types**

There are two primary types of provisioners:

<b>Types of Provisioners</b>	<b>Description</b>
Creation-Time Provisioner	Creation-time provisioners are only run during creation, not during updating or any other lifecycle  If a creation-time provisioner fails, the resource is marked as tainted.
Destroy-Time Provisioner	Destroy provisioners are run before the resource is destroyed.

If when = destroy is specified, the provisioner will run when the resource it is defined within is destroyed.

```

resource "aws_instance" "web" {
    # ...

    provisioner "local-exec" {
        when      = destroy
        command   = "echo 'Destroy-time provisioner'"
    }
}

```

## Module 4: Failure Behavior

By default, provisioners that fail will also cause the terraform apply itself to fail.

The `on_failure` setting can be used to change this. The allowed values are:

Allowed Values	Description
<code>continue</code>	Ignore the error and continue with creation or destruction.
<code>fail</code>	Raise an error and stop applying (the default behavior). If this is a creation provisioner, taint the resource.

```

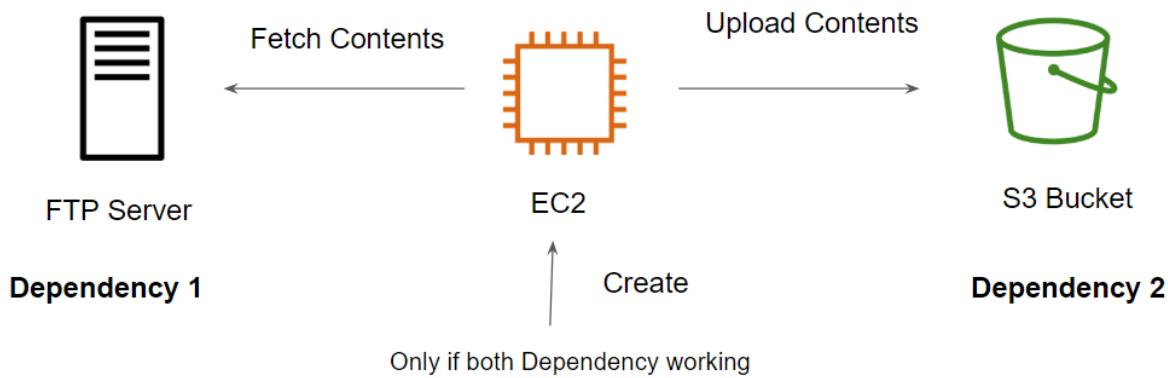
resource "aws_instance" "web" {
    # ...

    provisioner "local-exec" {
        command   = "echo The server's IP address is ${self.private_ip}"
        on_failure = continue
    }
}

```

## Module 5: Null Resource

The null\_resource implements the standard resource lifecycle but takes no further action.

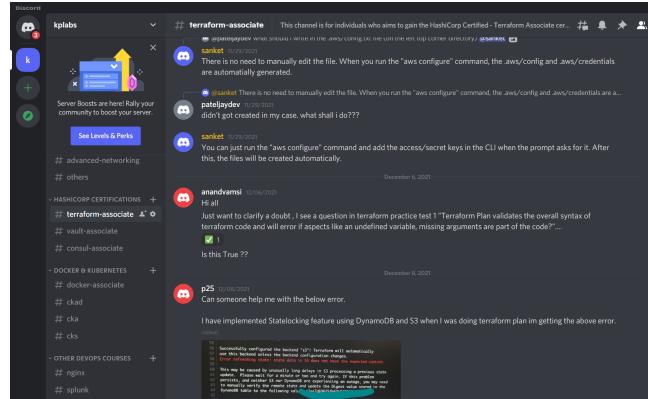


## Join Our Discord Community

We invite you to join our Discord community, where you can interact with our support team for any course-based technical queries and connect with other students who are doing the same course.

Joining URL:

<http://kplabs.in/chat>





# KPLABS Course

HashiCorp Certified: Terraform Associate

## Domain 4

**ISSUED BY**

Zeal

**REPRESENTATIVE**

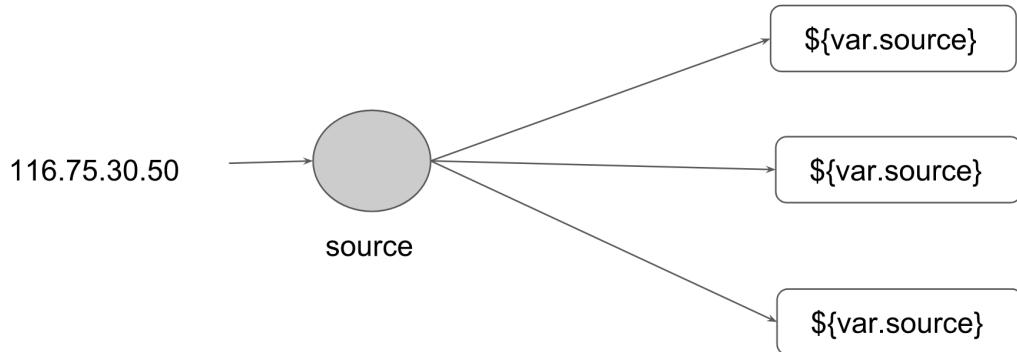
[instructors@kplabs.in](mailto:instructors@kplabs.in)

# Domain 4 - Terraform Modules & Workspaces

## Module 1: Understanding the DRY principle

In software engineering, don't repeat yourself (DRY) is a principle of software development aimed at reducing repetition of software patterns.

In the earlier lecture, we were making static content into variables so that there can be a single source of information.



### 1.1 Generic Scenario:

We do repeat multiple times various terraform resources for multiple projects.

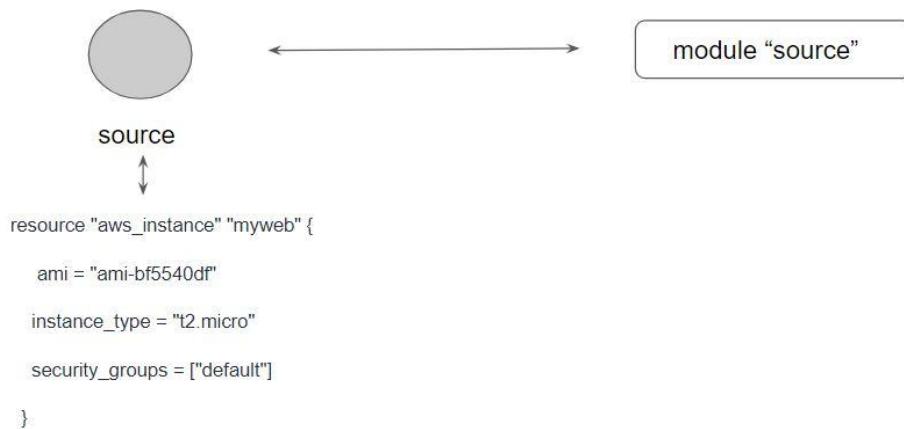
Sample EC2 Resource

```
resource "aws_instance" "myweb" {  
    ami = "ami-bf5540df"  
    instance_type = "t2.micro"  
    security_groups = ["default"]  
}
```

Instead of repeating a resource block multiple times, we can make use of a centralized structure.

### 1.2 Centralized Structure:

We can centralize the terraform resources and can call out from TF files whenever required.



## Module 2: Challenges with Terraform Modules

One common need for infrastructure management is to build environments like staging, production with a similar setup but keeping environment variables different.

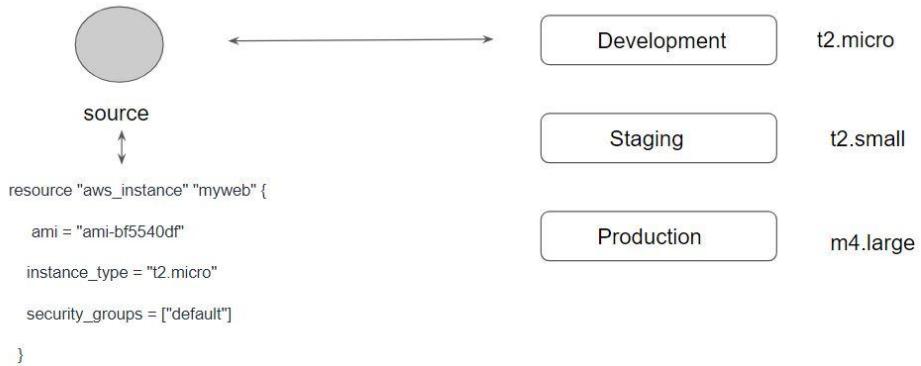
Staging

instance\_type = t2.micro

Production

instance\_type = m4.large

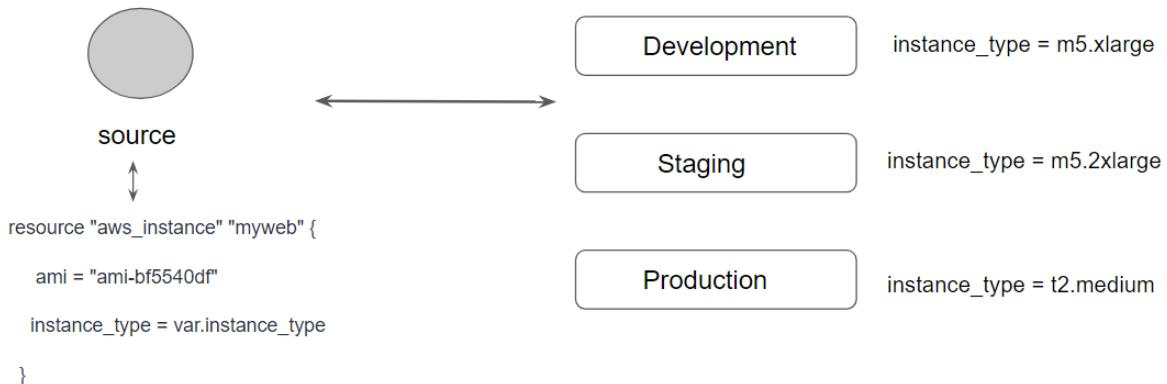
When we use modules directly, the resources will be a replica of code in the module.



## Module 3: Using Locals with Modules

### 3.1 Understanding the Challenge

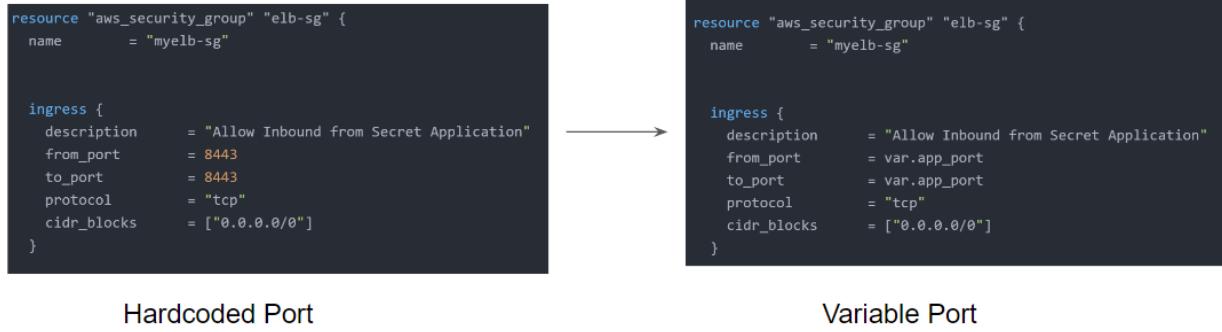
Using variables in Modules can also allow users to override the values which you might not want.



### 3.2 Setting the Context

There can be many repetitive values in modules and this can make your code difficult to maintain.

You can centralize these using variables but users will be able to override it.



### 3.3 Using Locals

Instead of variables, you can make use of locals to assign the values.

You can centralize these using variables but users will be able to override it.

```

resource "aws_security_group" "ec2-sg" {
  name      = "myec2-sg"

  ingress {
    description      = "Allow Inbound from Secret Application"
    from_port        = local.app_port
    to_port          = local.app_port
    protocol         = "tcp"
    cidr_blocks     = ["0.0.0.0/0"]
  }

  locals {
    app_port = 8443
  }
}

```

## Module 4: Module Outputs

### 4.1 Revising Output Values

Output values make information about your infrastructure available on the command line, and can expose information for other Terraform configurations to use.

```
output "instance_ip_addr" {
  value = aws_instance.server.private_ip
}
```

### 4.2 Accessing Child Module Outputs

In a parent module, outputs of child modules are available in expressions as `module.<MODULE NAME>.<OUTPUT NAME>`

```
resource "aws_security_group" "ec2-sg" {
  name      = "myec2-sg"

  ingress {
    description      = "Allow Inbound from Secret Application"
    from_port        = 8433
    to_port          = 8433
    protocol         = "tcp"
    cidr_blocks     = ["0.0.0.0/0"]
  }
}

output "sg_id" {
  value = aws_security_group.ec2-sg.arn
}
```

```
module "sgmodule" {
  source = "../../modules/sg"
}

resource "aws_instance" "web" {
  ami           = "ami-0ca285d4c2cda3300"
  instance_type = "t3.micro"
  vpc_security_group_ids = [module.sgmodule.sg_id]
}
```

## Module 5: Terraform Registry

The Terraform Registry is a repository of modules written by the Terraform community.

The registry can help you get started with Terraform more quickly



## 5.1 Module Location

If we intend to use a module, we need to define the path where the module files are present.

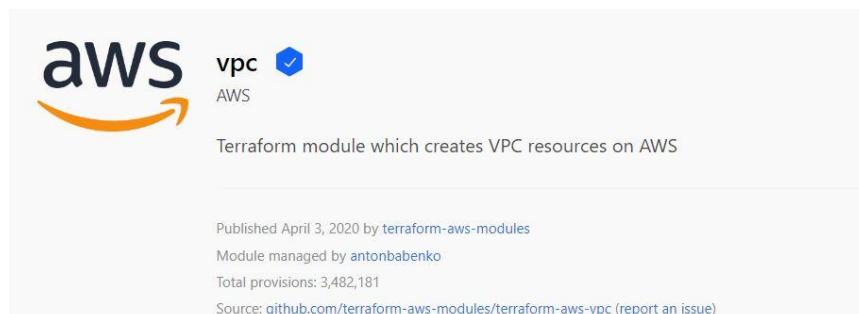
The module files can be stored in multiple locations, some of these include:

- Local Path
- GitHub
- Terraform Registry
- S3 Bucket
- HTTP URLs

## 5.2 Verified Modules in Terraform Registry

Within Terraform Registry, you can find verified modules that are maintained by various third-party vendors.

These modules are available for various resources like AWS VPC, RDS, ELB, and others



Verified modules are reviewed by HashiCorp and actively maintained by contributors to stay up-to-date and compatible with both Terraform and their respective providers.

The blue verification badge appears next to modules that are verified.

Module verification is currently a manual process restricted to a small group of trusted HashiCorp partners.

### 5.3 Using Registry Modules in Terraform

To use the Terraform Registry module within the code, we can make use of the source argument that contains the module path.

Below code references the EC2 Instance module within terraform registry.

```
module "ec2-instance" {  
  source = "terraform-aws-modules/ec2-instance/aws"  
  version = "2.13.0"  
  # insert the 10 required variables here  
}
```

## Module 6: Requirement for Publishing Modules in Terraform Registry

### 6.1 Overview of Publishing Modules

Anyone can publish and share modules on the Terraform Registry.

Published modules support versioning, automatically generate documentation, allow browsing version histories, show examples and READMEs, and more.

 <b>gruntwork-io / gke</b> Terraform code and scripts for deploying a Google Kubernetes Engine (GKE) cluster.  <span>🕒 9 months ago</span> <span>⬇️ 10.6K</span> <span> provider</span>	 <b>hashicorp / consul</b> A Terraform Module for how to run Consul on Google Cloud using Terraform and Packer.  <span>🕒 2 years ago</span> <span>⬇️ 6.6K</span> <span> provider</span>
 <b>gruntwork-io / sql</b> Terraform modules for deploying Google Cloud SQL (e.g. MySQL, PostgreSQL) in GCP.  <span>🕒 9 months ago</span> <span>⬇️ 4.3K</span> <span> provider</span>	 <b>gruntwork-io / static-assets</b> Modules for managing static assets (CSS, JS, images) in GCP.  <span>🕒 9 months ago</span> <span>⬇️ 22.4K</span> <span> provider</span>

## 6.2 Requirements for Publishing Module

Requirement	Description
GitHub	The module must be on GitHub and must be a public repo. This is only a requirement for the public registry.
Named	Module repositories must use this three-part name format terraform-<PROVIDER>-<NAME>
Repository description	The GitHub repository description is used to populate the short description of the module.
Standard module structure	The module must adhere to the standard module structure.
x.y.z tags for releases	The registry uses tags to identify module versions. Release tag names must be a semantic version, which can optionally be prefixed with a v. For example, v1.0.4 and 0.9.2

## 6.3 Standard Module Structure

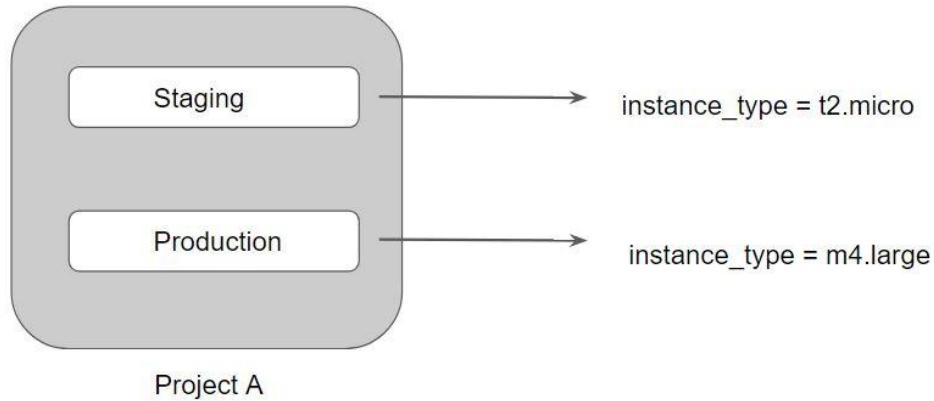
The standard module structure is a file and directory layout that is recommended for reusable modules distributed in separate repositories

```
$ tree minimal-module/
.
├── README.md
└── main.tf
└── variables.tf
└── outputs.tf
```

```
$ tree complete-module/
.
├── README.md
└── main.tf
└── variables.tf
└── outputs.tf
└── ...
└── modules/
    ├── nestedA/
    │   ├── README.md
    │   ├── variables.tf
    │   ├── main.tf
    │   └── outputs.tf
    └── nestedB/
        └── ...
└── examples/
    ├── exampleA/
    │   └── main.tf
    └── exampleB/
        └── .../
```

## Module 7: Terraform Workspace

Terraform allows us to have multiple workspaces, with each of the workspaces we can have a different set of environment variables associated



Terraform starts with a single workspace named "default".

This workspace is special both because it is the default and also because it cannot ever be deleted.

If you've never explicitly used workspaces, then you've only ever worked on the "default" workspace.

Workspaces are managed with the terraform workspace set of commands.

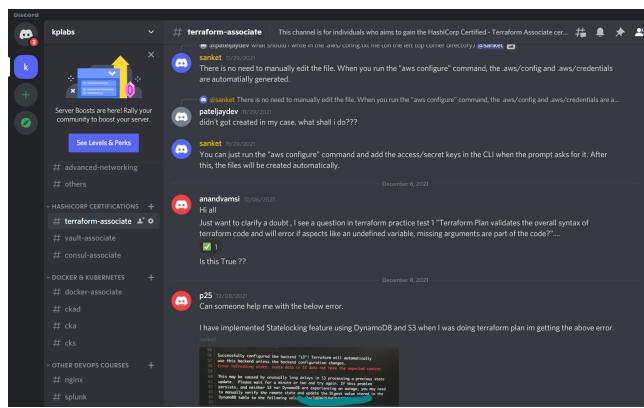
To create a new workspace and switch to it, you can use `terraform workspace new`; to switch workspaces you can use `terraform workspace select`; etc.

## Join Our Discord Community

We invite you to join our Discord community, where you can interact with our support team for any course-based technical queries and connect with other students who are doing the same course.

Joining URL:

<http://kplabs.in/chat>





# KPLABS Course

HashiCorp Certified: Terraform Associate

## Domain 5 - Remote State Management

**ISSUED BY**

Zeal

**REPRESENTATIVE**

[instructors@kplabs.in](mailto:instructors@kplabs.in)

# Domain 5 - Remote State Management

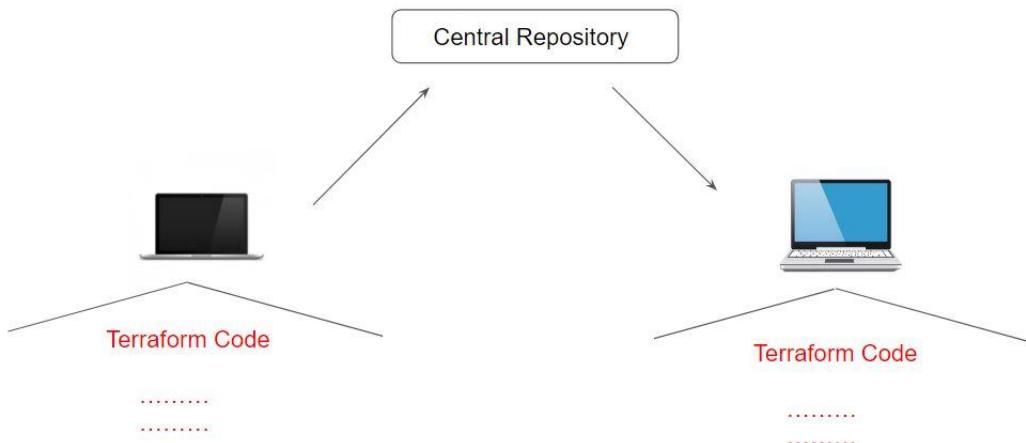
## Module 1: Integrating with GIT for team management

Till now, we have been working with terraform code locally.



However, storing your configuration files locally is not always an idea specifically in the scenario where other members of the team are also working on Terraform.

For such cases, it is important to store your Terraform code to a centralized repository like in Git.



## Module 2: Module Sources in Terraform

The source argument in a module block tells Terraform where to find the source code for the desired child module.

- Local paths
- Terraform Registry
- GitHub
- Bitbucket
- Generic Git, Mercurial repositories
- HTTP URLs
- S3 buckets
- GCS buckets

Let us explore some of the supported module sources.

### 2.1 Local Path

A local path must begin with either ./ or ../ to indicate that a local path is intended.

```
module "consul" {  
    source = "../consul"  
}
```

### 2.2 Git Module Source

Arbitrary Git repositories can be used by prefixing the address with the special git:: prefix.

After this prefix, any valid Git URL can be specified to select one of the protocols supported by Git.

```
module "vpc" {
  source = "git::https://example.com/vpc.git"
}

module "storage" {
  source = "git::ssh://username@example.com/storage.git"
}
```

## 2.3 Referencing to a Branch

By default, Terraform will clone and use the default branch (referenced by HEAD) in the selected repository.

You can override this using the ref argument:

```
module "vpc" {
  source = "git::https://example.com/vpc.git?ref=v1.2.0"
}
```

The value of the ref argument can be any reference that would be accepted by the git checkout command, including branch and tag names.

## Module 3: Terraform & GitIgnore

The .gitignore file is a text file that tells Git which files or folders to ignore in a project.

.gitignore
conf/
*.artifacts
credentials



Depending on the environment, it is recommended to avoid committing certain files to GIT.

Files to Ignore	Description
.terraform	This file will be recreated when terraform init is run.
terraform.tfvars	Likely to contain sensitive data like usernames/passwords and secrets.
terraform.tfstate	Should be stored in the remote side.
crash.log	If terraform crashes, the logs are stored to a file named crash.log

## Module 4: Terraform Backend

### 4.1 Basics of Backends

Backends primarily determine where Terraform stores its state.

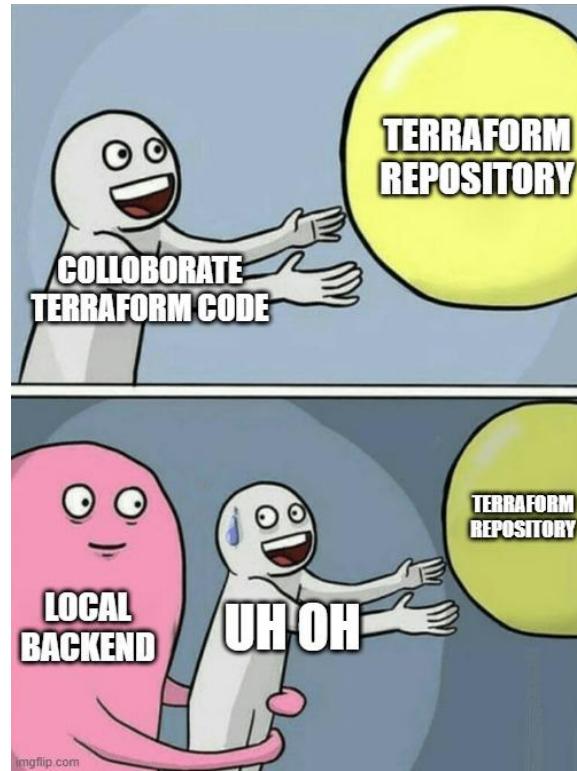
By default, Terraform implicitly uses a backend called local to store state as a local file on disk.



### 4.2 Challenge with Local Backend

Nowadays Terraform projects are handled and collaborated by an entire team.

Storing the state file in the local laptop will not allow collaboration.

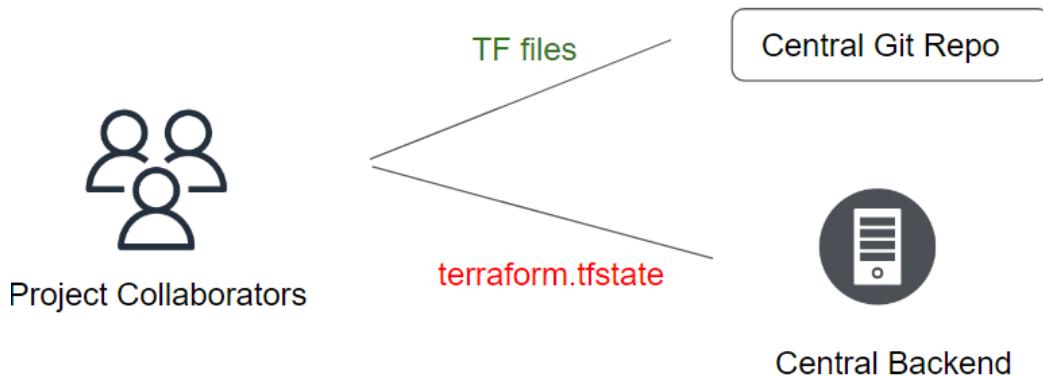


#### 4.3 Ideal Architecture

Following describes one of the recommended architectures:

The Terraform Code is stored in Git Repository.

The State file is stored in a Central backend.



#### 4.4 Backends Supported in Terraform

Terraform supports multiple backends that allow remote service related operations.

Some of the popular backends include:

- S3
- Consul
- Azurerm
- Kubernetes
- HTTP
- ETCD

#### 4.5 Important Note

Accessing state in a remote service generally requires some kind of access credentials

Some backends act like plain "remote disks" for state files; others support locking the state while operations are being performed, which helps prevent conflicts and inconsistencies.



## Module 5: State File Locking

### 5.1 Understanding State Lock

Whenever you are performing write operation, terraform would lock the state file.

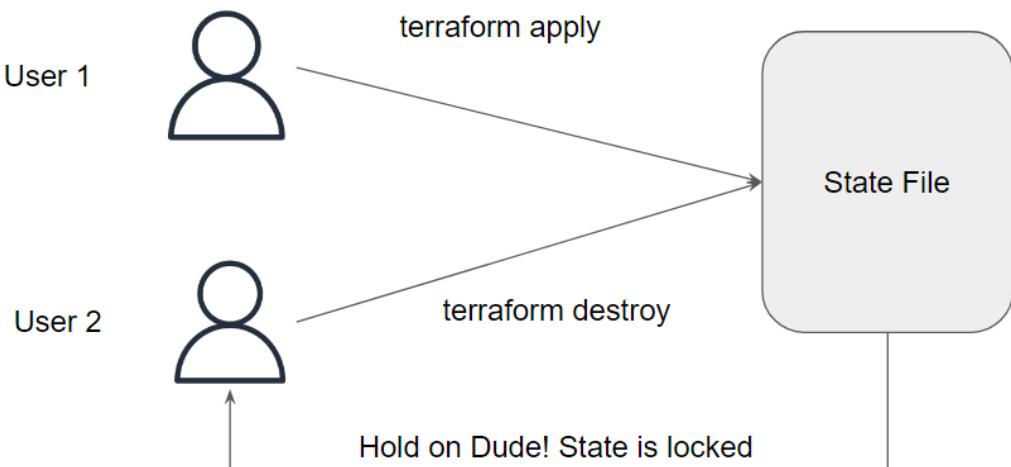
This is very important as otherwise during your ongoing terraform apply operations, if others also try for the same, it can corrupt your state file.

```
C:\Users\Zeal Vora\Desktop\tf-demo\remote-backend>terraform plan
Error: Error acquiring the state lock

Error message: Failed to read state file: The state file could not be read: read terraform.tfstate: The process
cannot access the file because another process has locked a portion of the file.

Terraform acquires a state lock to protect the state from being written
by multiple users at the same time. Please resolve the issue above and try
again. For most commands, you can disable locking with the "-lock=false"
flag, but this is not recommended.
```

Following diagram illustrates the basic working.



## 5.2 Important Note

State locking happens automatically on all operations that could write state. You won't see any message that it is happening

If state locking fails, Terraform will not continue

Not all backends support locking. The documentation for each backend includes details on whether it supports locking or not.

## 5.3 Force Unlocking State

Terraform has a **force-unlock** command to manually unlock the state if unlocking failed.

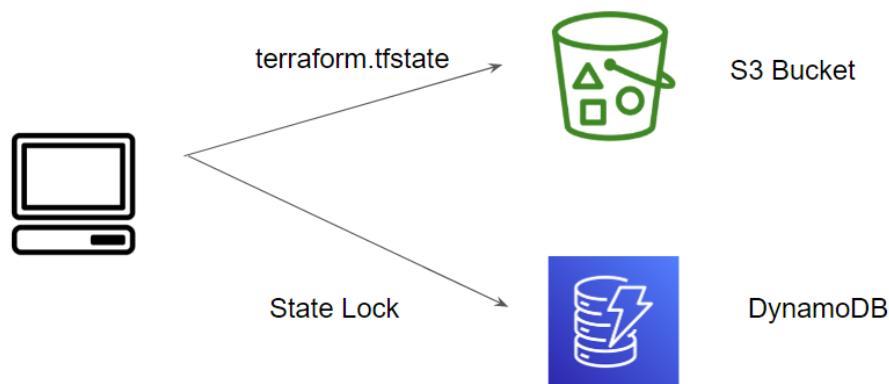
If you unlock the state when someone else is holding the lock it could cause multiple writers.

Force unlock should only be used to unlock your own lock in the situation where automatic unlocking failed.

## Module 6: State Locking in S3 Backend

By default, S3 does not support State Locking functionality.

You need to make use of the DynamoDB table to achieve state locking functionality.



## Module 7: Terraform State Management

As your Terraform usage becomes more advanced, there are some cases where you may need to modify the Terraform state.

It is important to never modify the state file directly. Instead, make use of `terraform state` command.

There are multiple sub-commands that can be used with `terraform state`, these include:

State Sub Command	Description
list	List resources within terraform state file.
mv	Moves item with terraform state.
pull	Manually download and output the state from remote state.
push	Manually upload a local state file to remote state.
rm	Remove items from the Terraform state
show	Show the attributes of a single resource in the state.

### 7.1 Sub Command - List

The `terraform state list` command is used to list resources within a Terraform state.

```
bash-4.2# terraform state list
aws_iam_user.lb
aws_instance.webapp
```

### 7.2 Sub Command - Move

The `terraform state mv` command is used to move items in a Terraform state.

This command is used in many cases in which you want to rename an existing resource without destroying and recreating it.

Due to the destructive nature of this command, this command will output a backup copy of the state prior to saving any changes

Overall Syntax:

```
terraform state mv [options] SOURCE DESTINATION
```

### 7.3 Sub Command - Pull

The terraform state pull command is used to manually download and output the state from a remote state.

This is useful for reading values out of state (potentially pairing this command with something like jq).

### 7.4 Sub Command - Push

The terraform state push command is used to manually upload a local state file to remote state.

This command should rarely be used.

### 7.5 Sub Command - Remove

The terraform state rm command is used to remove items from the Terraform state.

Items removed from the Terraform state are not physically destroyed.

Items removed from the Terraform state are only no longer managed by Terraform

For example, if you remove an AWS instance from the state, the AWS instance will continue running, but terraform plan will no longer see that instance.

## 7.6 Sub Command - Show

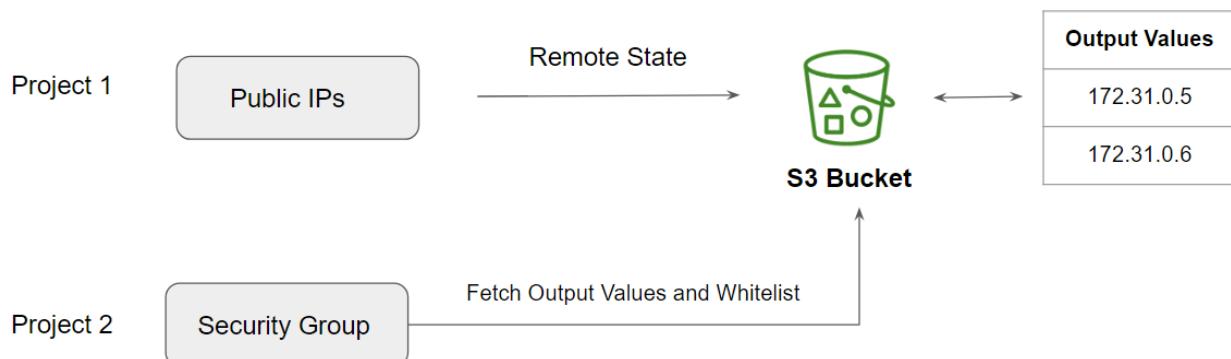
The terraform state show command is used to show the attributes of a single resource in the Terraform state.

```
bash-4.2# terraform state show aws_instance.webapp
# aws_instance.webapp:
resource "aws_instance" "webapp" {
  ami                      = "ami-082b5a644766e0e6f"
  arn                     = "arn:aws:ec2:us-west-2:018721151861:instance/i-0107ea9ed06c467e0"
  associate_public_ip_address = true
  availability_zone        = "us-west-2b"
  cpu_core_count           = 1
  cpu_threads_per_core     = 1
  disable_api_termination  = false
  ebs_optimized            = false
  get_password_data        = false
  id                       = "i-0107ea9ed06c467e0"
  instance_state            = "running"
  instance_type             = "t2.micro"
```

## Module 8: Connecting Remote States

### 8.1 Basics of Terraform Remote State

The [terraform\\_remote\\_state](#) data source retrieves the root module output values from some other Terraform configuration, using the latest state snapshot from the remote backend.



## 8.2 Step 1 - Create a Project with Output Values & S3 Backend

```
resource "aws_eip" "lb" {
  vpc      = true
}

output "eip_addr" {
  value = aws_eip.lb.public_ip
}
```



```
terraform {
  backend "s3" {
    bucket = "kplabs-terraform-backend"
    key    = "network/eip.tfstate"
    region = "us-east-1"
  }
}
```

## 8.3 Step 2 - Reference Output Values from Different Project

```
data "terraform_remote_state" "eip" {
  backend = "s3"
  config = {
    bucket = "kplabs-terraform-backend"
    key    = "network/eip.tfstate"
    region = "us-east-1"
  }
}
```



```
resource "aws_security_group" "allow_tls" {
  name      = "allow_tls"
  description = "Allow TLS inbound traffic"

  ingress {
    description     = "TLS from VPC"
    from_port       = 443
    to_port         = 443
    protocol        = "tcp"
    cidr_blocks    = ["${data.terraform_remote_state.eip.outputs.eip_addr}/32"]
  }
}
```

## Module 9: Terraform Import

It might happen that there is a resource that is already created manually.

In such a case, any change you want to make to that resource must be done manually.



Terraform is able to import existing infrastructure. This allows you to take resources you've created by some other means and bring it under Terraform management.

The current implementation of Terraform import can only import resources into the state. It does not generate configuration. A future version of Terraform will also generate configuration.

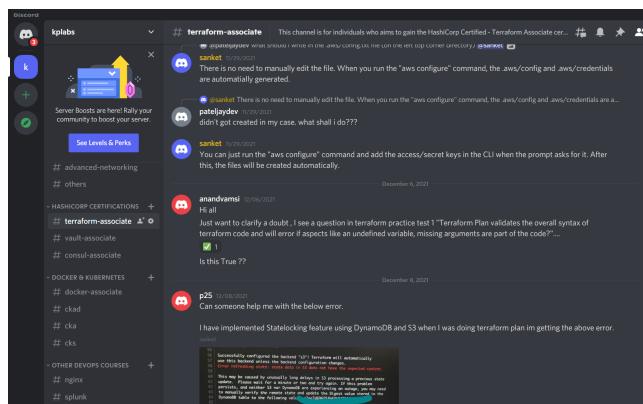
Because of this, prior to running terraform import it is necessary to write manually a resource configuration block for the resource, to which the imported object will be mapped.

## Join Our Discord Community

We invite you to join our Discord community, where you can interact with our support team for any course-based technical queries and connect with other students who are doing the same course.

Joining URL:

<http://kplabs.in/chat>





# KPLABS Course

HashiCorp Certified: Terraform Associate

## Domain 6

**ISSUED BY**

Zeal

**REPRESENTATIVE**

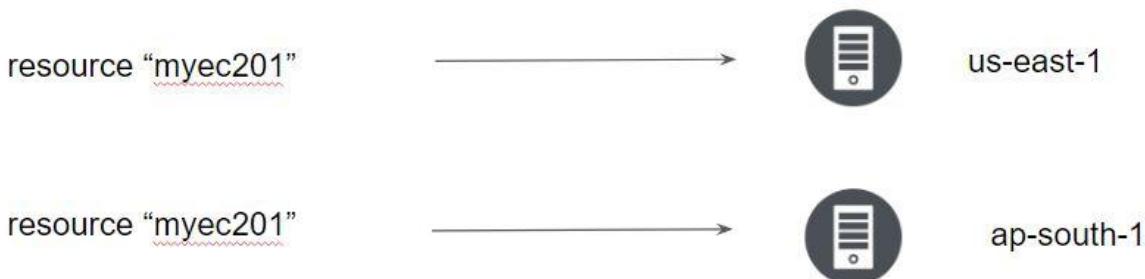
[instructors@kplabs.in](mailto:instructors@kplabs.in)

# Domain 6 - Security Primer

## Module 1: Provider Configuration

Till now, we have been hardcoding the aws-region parameter within the providers.tf

This means that resources would be created in the region specified in the providers.tf file.



By default, resources use a default provider configuration inferred from the first word of the resource type name.

For example, a resource of type `aws_instance` uses the default (un-aliased) `aws` provider configuration unless otherwise stated.

To select an aliased provider for a resource or data source, set its provider meta-argument to a `<PROVIDER NAME>.<ALIAS>` reference:

```
resource "aws_instance" "foo" {
  provider = aws.west

  #
}
```

## Module 2: Handling Multiple AWS Profiles in Terraform

You can optionally define multiple configurations for the same provider, and select which one to use on a per-resource or per-module basis.

The primary reason for this is to support multiple regions for a cloud platform.

To include multiple configurations for a given provider, include multiple provider blocks with the same provider name, but set the alias meta-argument to an alias name to use for each additional configuration. For example:

```
# The default provider configuration
provider "aws" {
  region = "us-east-1"
}

# Additional provider configuration for west coast region
provider "aws" {
  alias  = "west"
  region = "us-west-2"
}
```

The provider block without alias set is known as the default provider configuration.

When an alias is set, it creates an additional provider configuration.

For providers that have no required configuration arguments, the implied empty configuration is considered to be the default provider configuration.

## Module 3: Sensitive Parameter

With the organization managing its entire infrastructure in terraform, it is likely that you will see some sensitive information embedded in the code.

When working with a field that contains information likely to be considered sensitive, it is best to set the Sensitive property on its schema to true

```
output "db_password" {
  value      = aws_db_instance.db.password
  description = "The password for logging in to the database."
  sensitive   = true
}
```

Setting the sensitive to “true” will prevent the field's values from showing up in CLI output and in Terraform Cloud

It will not encrypt or obscure the value in the state, however.

```
C:\Users\Zeal Vora\Desktop\terraform\sensitive data>terraform apply

Apply complete! Resources: 0 added, 0 changed, 0 destroyed.

Outputs:

db_password = <sensitive>
```

## Module 4: Overview of HashiCorp Vault

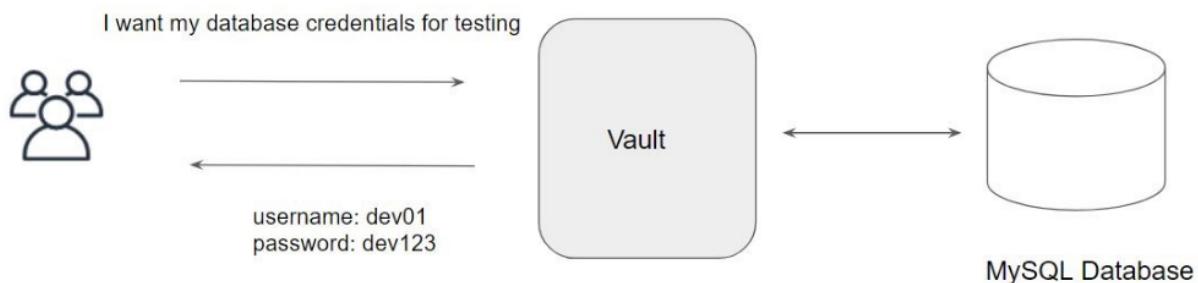
HashiCorp Vault allows organizations to securely store secrets like tokens, passwords, certificates along with access management for protecting secrets.

One of the common challenges nowadays in an organization is “Secrets Management”

Secrets can include database passwords, AWS access/secret keys, API Tokens, encryption keys and others.

The screenshot shows the HashiCorp Vault web interface. At the top, there are navigation tabs: Secrets (which is selected), Access, Policies, and Tools. On the right side of the header, there are Status, Help, and Settings icons. Below the header, the main content area is titled "Secrets Engines". It lists several engines: "aws aws/" (selected, indicated by a cursor icon over it), "cubbyhole cubbyhole\_b18723c3", "secret secret\_kv\_ccca51e2", "transit transit\_05f7d94c", and "consul consul\_a31f034b". Each engine entry has a three-dot menu icon on its right. At the top right of the "Secrets Engines" section, there is a link "Enable new engine >".

Vault can also generate dynamic secrets like AWS Access/Secret keys that has validity for a limited time.

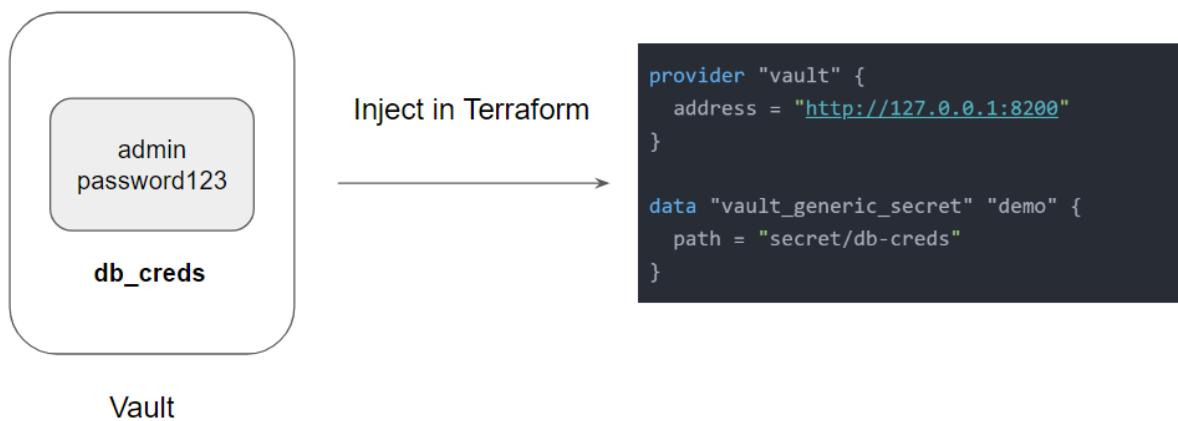


Once Vault is integrated with multiple backends, your life will become much easier and you can focus more on the right work.

Major aspects related to Access Management can be taken over by vault.

## Module 5: Terraform and Vault Integration

The Vault provider allows Terraform to read from, write to, and configure HashiCorp Vault.



### Important Note:

Interacting with Vault from Terraform causes any secrets that you read and write to be persisted in Terraform's state file.





# KPLABS Course

HashiCorp Certified: Terraform Associate

## Domain 7

**ISSUED BY**

Zeal

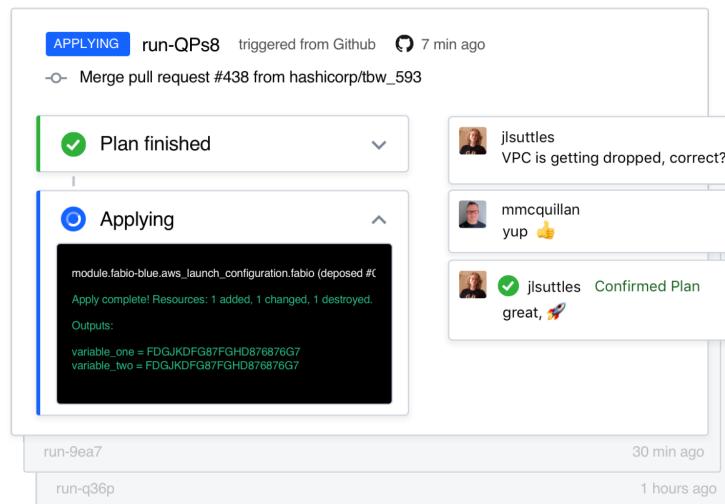
**REPRESENTATIVE**

[instructors@kplabs.in](mailto:instructors@kplabs.in)

# Domain 7 - Terraform Cloud & Enterprise Capabilities

## Module 1: Overview of Terraform Cloud

Terraform Cloud manages Terraform runs in a consistent and reliable environment with various features like access controls, private registry for sharing modules, policy controls, and others.



Terraform Cloud is available as a hosted service at <https://app.terraform.io>.

## Module 2: Overview of Sentinel

Sentinel is an embedded policy-as-code framework integrated with the HashiCorp Enterprise products.

It enables fine-grained, logic-based policy decisions, and can be extended to use information from external sources.

Note: Sentinel policies are paid feature



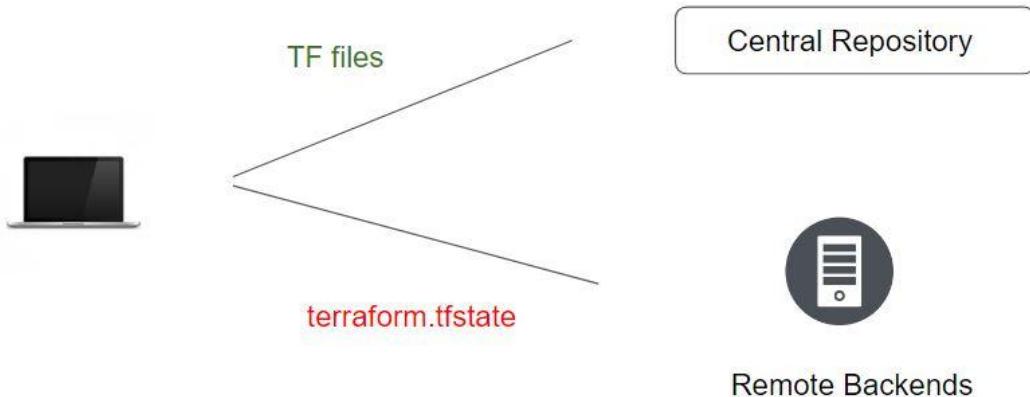
## Module 3: Remote Backends

Terraform supports various types of remote backends which can be used to store state data.

As of now, we were storing state data in local and GIT repository.

Depending on remote backends that are being used, there can be various features.

- Standard BackEnd Type: State Storage and Locking
- Enhanced BackEnd Type: All features of Standard + Remote Management



When using full remote operations, operations like `terraform plan` or `terraform apply` can be executed in Terraform Cloud's run environment, with log output streaming to the local terminal.

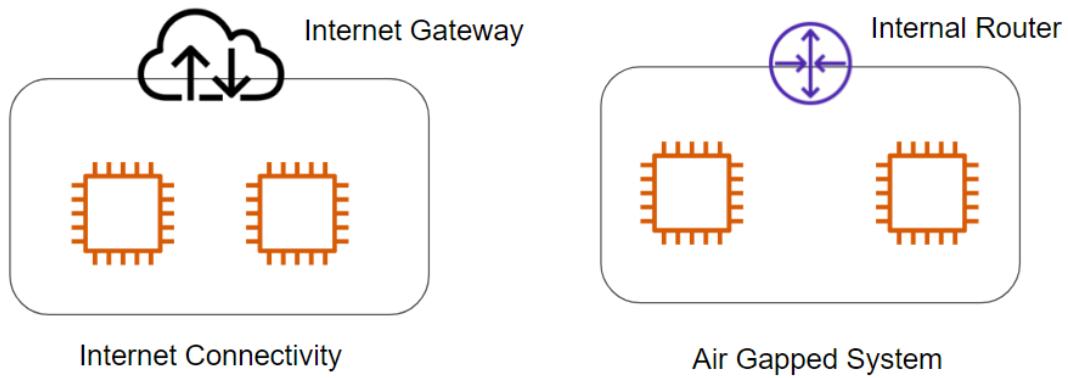
Remote plans and applies use variable values from the associated Terraform Cloud workspace.

Terraform Cloud can also be used with local operations, in which case only state is stored in the Terraform Cloud backend.

## Module 4: Air Gapped Environments

### 4.1 Understanding Concept of Air Gap

An air gap is a network security measure employed to ensure that a secure computer network is physically isolated from unsecured networks, such as the public Internet.



### 4.2 Usage of Air Gapped Systems

Air Gapped Environments are used in various areas. Some of these include:

- Military/governmental computer networks/systems
- Financial computer systems, such as stock exchanges
- Industrial control systems, such as SCADA in Oil & Gas fields

### 4.3 Terraform Enterprise Installation Methods

Terraform Enterprise installs using either an online or air gapped method and as the names infer, one requires internet connectivity, the other does not

WORKSPACE NAME	RUN STATUS	LATEST CHANGE	RUN	REPO
exoscale-test	✓ <span>APPLIED</span>	5 months ago	run-884c	NICK/tf-exoscale-minimum
tf-test-dev	✗ <span>FAILED</span>	3 months ago	run-955c	nagerlund/terraform-tf-test
migrated-test	✓ <span>PLANNED</span>	5 months ago	run-895f	nagerlund/terraform-minimum
migrated-first	✓ <span>PLANNED</span>	5 months ago	run-924p	nagerlund/terraform-minimum
migrated-second	✓ <span>PLANNED</span>	5 months ago	run-894v	nagerlund/terraform-minimum
migrated-third	✓ <span>APPLIED</span>	5 months ago	run-988k	NICK/tf-exoscale-minimum
migrated-fourth	✓ <span>PLANNED</span>	5 months ago	run-8837	nagerlund/terraform-minimum
migrate-first-2	✗ <span>NEEDS CONFIRMATION</span>	3 months ago	run-9827	nagerlund/terraform-minimum

Terraform Enterprise

## Air Gap Install



Isolated Server

Choose your installation type



Please choose an installation type to continue.

[Continue »](#)

## Provide path or upload airgap bundle

Provide absolute path on this server to archive file

e.g. /mnt/installers/package.airgap

[Continue »](#)

Select file for upload

 [Upload Airgap Bundle](#)

To upload an app bundle, file must have a `.airgap` extension.

[« Back](#)